
Lecture 7: RTL Coding Guidelines

Slides adapted from Ofer Shacham
with contributions from Zain Asgar

Overview

Reading


- + RTL Coding Styles That Yield Simulation and Synthesis Mismatches:
https://www.synopsys.com/news/pubs/eurosnug/eurosnug2001/mills_rtl_coding_sfinal.pdf
- + Predicting Routability at the Register Transfer Level:
http://www.synopsys.com/news/pubs/sjsnug/sj03/marshall_pres.pdf
- + RTL Coding & Optimization Guide:
http://www.synopsys.com/news/pubs/sjsnug/sj02/marshall_slidesf.pdf
- + Retiming:
<http://www.eecs.berkeley.edu/~keutzer/classes/244fa2005/lectures/8-2-retiming.pdf>
- + “System Verilog For Design,” Sutherland, Davidmann & Flake

Background

While we assume you have all coded in Verilog before, if you want to create a complex system, you need to stylize the type of Verilog you create. We will talk about how to do a little preplanning, some general style rules, leveraging retiming to do some of your job, and then how to avoid / deal with communication issues.

Why Use Coding Guidelines: The Human Perspective

- Coding guidelines make the chip design process more efficient
 - Today's designs are thousands of verilog lines
 - Written by tens/hundreds of people
 - Teams from all around the world design one chip
- Coding guidelines are set to save us from ourselves
 - Think about reading code one year after the project is done




Why Use Coding Guidelines: The Tools Perspective

- Simulators can do magical things
 - But real hardware can't!

```
initial begin
  integer I;
  for (i=0; i<128; i++) begin
    mem[i] = 0;
  end
end
```

- Some issues may only be realized at the place & route phase



This Lecture is Based On...

- Lessons learned the hard way
 - Or the easy way (i.e. other people's hard way ;-))
- Lessons learned from Zain Asgar's presentation
- A collection of other people's presentations/papers/books
 - See the reading list for this lecture!

Planning Is Critical: Things to Do Before RTL

- Do you have a specification for your design?
- Create a block level drawing first
 - Break the design down to manageable chunks of RTL code
 - Look at the communication paths between the chunks
- Any pre-existing physical requirements for the design?
 - Is area a limitation? Is power or frequency a limitation?
 - Do you have enough room for all your I/Os?
- What is the expected frequency of operation? what technology?
 - Roughly estimate number of logic levels between flops
 - Determine at what level(s) to register outputs
 - Where are the anticipated critical paths?

For Large Designs, You Must: Decompose & Decouple

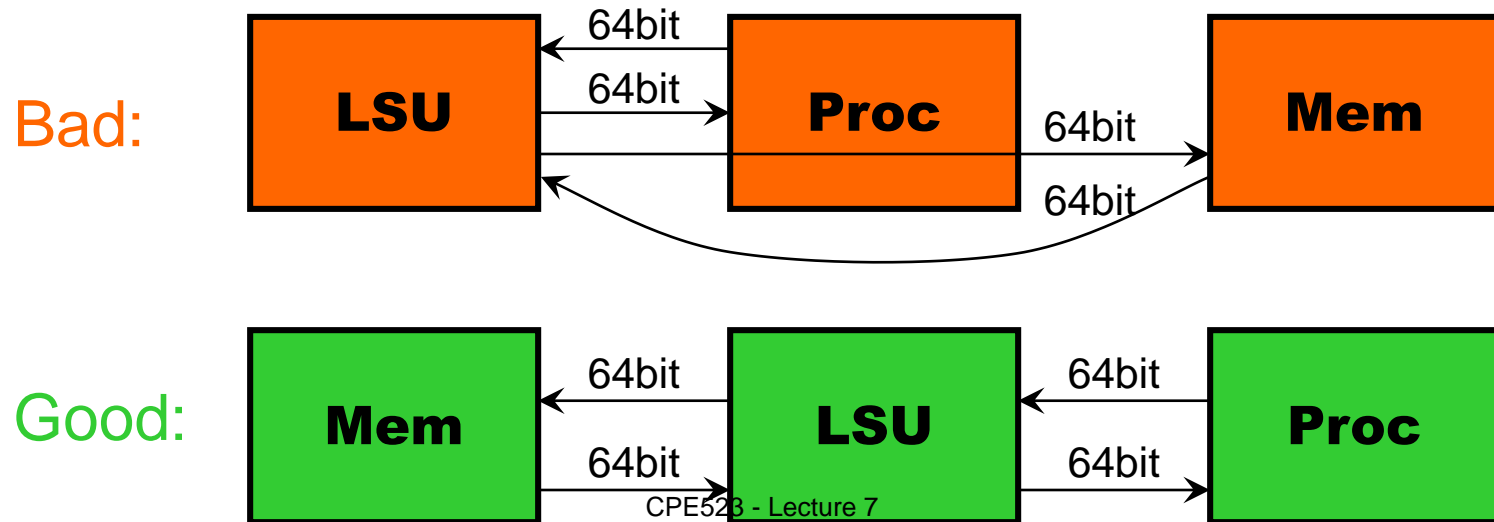
- Think about how the blocks interact
 - Decouple state machines dependencies between blocks
 - Separate data path logic and control logic
- Start thinking about physical constraints
 - What size can your synthesis/physical design tools handle?
 - Currently it is around 200K gates (and rising)
 - Wires will be pretty long between these (large) blocks
- At the top level - constrain blocks to around 200K gates
 - This will be the layout partition
 - Flop all outputs (no need to flop inputs)
 - Decouples timing between blocks (critical path is within a block)

Flopping Outputs


- Want to create a modular design
 - Need to be able to build / evaluate each block
- Flopping all outputs sets the required timing on all outputs
 - Must settle before setup time of flop
- Input can arrive from other units
 - Top level designs need to estimate this time
 - But the timing of inputs is just from wiring
 - Can easily estimate from a rough layout plan
- But this means a signal can only go through one block/cycle
 - In some cases the communication takes a full cycle
 - Then you flop input and output

Start Planning Wires

- Create a rough floorplan
 - Units that communicate should be adjacent (when possible)
 - Make sure that not too many wires cross or loop around
- Think carefully about nets with large fanouts.
 - These can have a huge timing disparity in synthesis vs P&R
 - They also make routing hard



One More Point About Wire Planning



- While we tend to draw wires that
 - Start at the edge of a block
 - End at the edge of the next block
- Wires origins and destinations are random in a block
- So for top level wire cap / delay estimates
 - Assume all wires start from middle of the physical block

Keep Clean Interfaces

- Keep interfaces simple
 - Don't overload signals unless absolutely necessary
 - Separate meaningful signals
- Key to successful **verification**, IP reuse, code robustness...

Bad:

```
input [73:0] proc2lsu; // interface processor to ld/st unit
output [89:0] lsu2proc;
assign proc2lsu = {req, op, addr, data...};
assign addr = lsu2proc[37:19];
```

Good:

```
input proc2lsu_req;
input [7:0] proc2lsu_op;
input [31:0] proc2lsu_addr;
input [48:0] proc2lsu_data;
```

Better: Use a SV interface

```
interface proc2lsu (input logic Clk);
    logic req;
    ...
endinterface
```

Files and Folders

- Separate design from environment files!
 - Even if the environment files are in Verilog / System Verilog
 - Separate the synthesis / place-and-route scripts
- It is conventional to name files **exactly** as the module inside
 - Some tools expect it

Bad:

```
// File name: Arithmetic_Logic.v  
module alu (...);  
...  

```

Good:

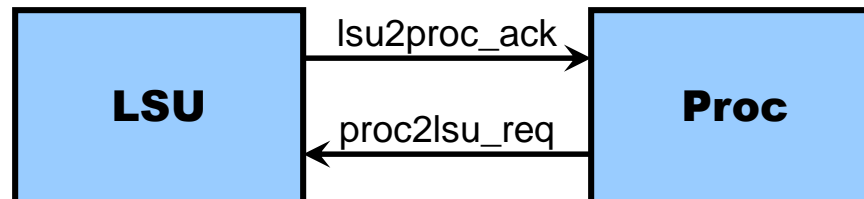
```
// File name: alu.v  
// Description: arithmetic-logic unit ...  
module alu (...);  

```

- For shared constants/parameters use headers (or SV packages)
 - Don't define twice (different namespace resolution between tools)

Naming Conventions

- Have a naming convention!
- Verilog is case sensitive
 - Avoid mixing CapitalStyle with underscore_style
- Avoid names that look/sound like keywords
 - Bad: “wire module_proc_input;”
- Interface signals
 - Don’t use the signal_in / signal_out – it’s confusing
 - I like the “src2dst_signal” notation



Naming Conventions, Cont'd

- Pervasive signals should be named the same
 - Or you end-up with {Clk, CK, CLK, clk} or {rst, Reset, reset}
- Some signals are commonly named by their function
 - Enable: signal_en
 - Negation: signal_b or signal_n
 - Delayed by one clock cycle: signal_d or signal_del
 - Request / Acknowledge / Grant: signal_req / _ack / _gnt
- Some signals are named by their qualifying clock
 - signal_clk1 is clocked by clk1
 - signal_clk2 is clocked by clk2

Verilog Constructs

- Verilog (and VHDL) sucks...
 - Makes it easy to write un-synthesizable code
 - Or code which is hard to implement
 - Or code that will behave differently before/after synthesis
- It is your job to understand the limitation
- It is your job to use the right construct for the right purpose
 - Use “programming” extensions where it makes sense
 - Verification
 - Only use synthesizable subset for your design

Verilog Constructs: Tasks and Functions

- Not the same as software methods/functions
 - Multiple function calls occur across die space, not time
- Don't use tasks
 - They only work well in simulation
- Be careful with functions
 - A function can only be a block of combinatorial logic and hence synthesizable
 - Sequential logic not allowed in functions
 - If your logic infers a latch → synthesis will not match simulation
 - Worse yet the synthesis tool will not complain

Verilog Constructs: Function Example


```
module bad_latch(  
    output o,  
    input a,  
    input en,  
    input nrst);  
  
    function latch;  
        input a, nrst, en;  
        if (!nrst) o = 1'b0;  
        else if (en) o = a;  
    endfunction  
  
    assign o = latch(a, nrst, en);  
endmodule
```

Synthesized result!



```
module good_latch(  
    output reg o,  
    input a,  
    input en,  
    input nrst);  
  
    always @(a or rst or en)  
        if (!nrst) o = 1'b0;  
        else if (en) o = a;  
endmodule
```

Synthesized result



Verilog Constructs: Loops

- Be careful with loops
 - They usually create more hardware
 - Make sure there is a constant termination
 - Loops with a variable termination are not synthesizable
- Rethink – do you really want a loop or a tree?
 - You are usually better off creating a tree structure yourself
 - The synthesis tool might try to create a tree for better timing
 - Don't make the synthesis tool do the designers job

Verilog Constructs: Loop Examples (That Give Correct Results In Simulation)

```
module multiplier(input [31:0] A, input [31:0] B, output [63:0] Res);  
    integer i;  
    logic [63:0] partial_prod[32]; // two dimensional array to hold partial products  
    logic [63:0] partial_sum[32]; // two dimensional array to hold partial sums  
    always_comb begin : Good_Loop  
        for (i=0; i<32; i=i+1)  
            partial_prod[i] = (A[i] == 1'b1)? { '0, B<<i} : '0;  
    end  
  
    always_comb begin : Bad_Loop  
        partial_sum[0] = partial_prod[0];  
        for (j=1; j<32; j++)  
            partial_sum[j] = partial_sum[j-1] + partial_prod[j];  
    end  
  
    assign Res = partial_sum[31];  
endmodule
```

This is parallel, independent logic

This creates a path of 32 adders!!

- How would you design the partial product summation?

Verilog Constructs: Always Blocks

- **Don't mix comb. and seq. logic in same always block!**
 - Logic is much clearer
 - Avoids weird issues with resets
 - Avoids latch/flop inferring that you did not plan

Bad1: mixing comb. and seq. logic

```
always @posedge Clk  
    counter <= counter + 1
```

Bad2: unintended latch inferred

```
always @ signal_a  
    if (signal_a)  
        signal_b = signal_a + 1;
```

Good1:

```
always @*  
    nx_counter = counter + 1
```

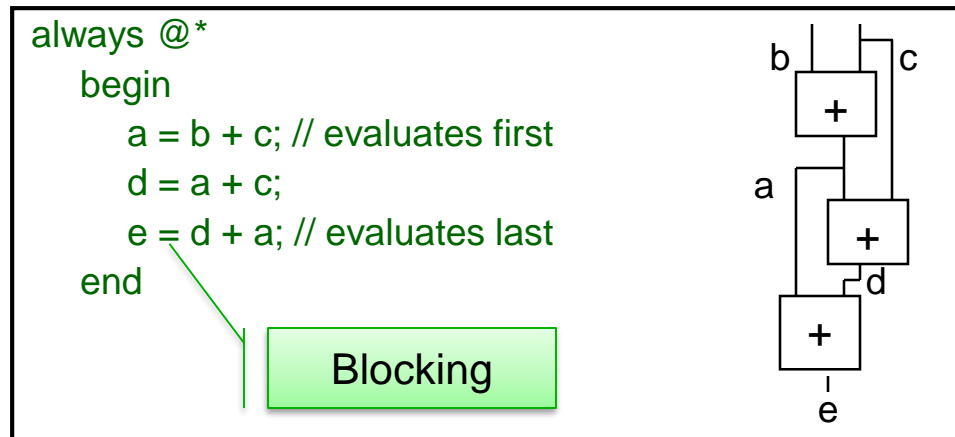
```
always @posedge Clk  
    counter <= nx_counter
```

Good2:

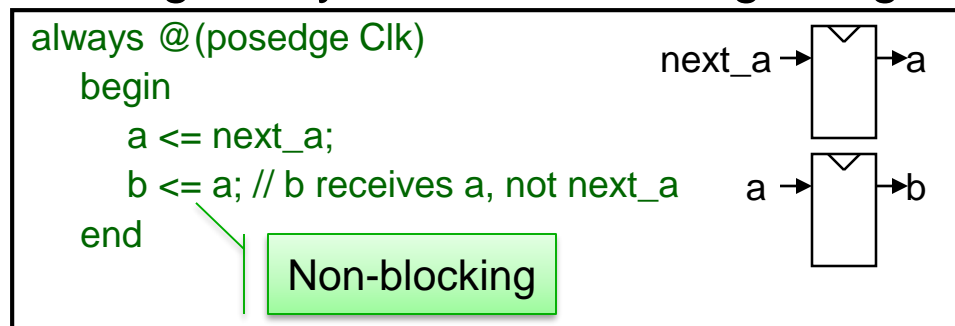
```
always @ signal_a  
    if (signal_a)  
        signal_b = signal_a + 1;  
    else  
        signal_b = signal_a + 2;
```

Verilog Constructs: Always Blocks

- For combinational logic only use blocking assignments



- For sequential logic only use non-blocking assignments



Verilog Constructs: Always Blocks In SV

- System Verilog attempts to make things clearer
 - Adds 'always_comb', 'always_latch' and 'always_ff'
 - Designer intent is clearer
- Note that (unfortunately) tools are not required to check your code
 - If you put combinational logic in an 'always_ff' you get weird results
 - It is good practice to synthesize and trace back any flop/latch

Verilog Constructs: Case Statements

- A *case* by definition is an “if... else if... else if... [else]”
 - This means that there is priority implicitly in the code
 - Yet, tools would try to optimize away priority if they can
- Always have defaults
 - Unless you meant to infer a latch
 - But then again, “**Don’t mix comb. and seq. logic!**”
- Try to fully enumerate all cases

Verilog Constructs: Case Pragmas

- In verilog you might see “// synopsys parallel_case full_case”
 - This is a **synthesis** pragma which **should be avoided**
 - parallel_case → synthesis tool should remove priority
 - full_case → synthesis tool should optimize outputs for unspecified cases
- System Verilog extends verilog: ‘unique case’ and ‘priority case’
 - ‘unique case’ is the safe way to do “// synopsys parallel_case full_case”
 - ‘priority case’ is the safe way to do “// synopsys full_case”
- ‘unique case’ tells the tools (both sim. and synth.) that:
 - One and only one *case_select* matches the *case_expr*.
- ‘priority case’ tells the tools (both sim. and synth.) that:
 - At least one *case_select* matches the *case_expr*.
 - Must be evaluated in order in which it was written

Verilog Constructs: Onehot Case

- One hot state machines are tricky
 - There is no way for the tool to know you wanted one hot
 - Some libraries have special (faster) one hot mux cells

The right way to specify a one-hot mux in system verilog:

```
reg [2:0] state_vec;
```

```
typedef enum { IDLE=0, NEXT, PREV} states_t;
```

Similar to verilog's 'parameter' or 'define'

```
always_ff @(posedge Clk)
```

```
if (Reset) state_vec <= 3'b1;
```

```
else state_vec <= next_state_vec;
```

Separated sequential logic from combinational logic

```
always_comb begin
```

```
next_state_vec = state_vec; // default
```

```
unique case (1'b1)
```

```
state_vec[IDLE]: ...
```

```
state_vec[NEXT]:...
```

```
state_vec[PREV]: ...
```

```
endcase
```

```
end
```

Instead of the SV 'unique case', you can use:
case (1'b1) '// synopsys parallel_case full_case'

- Always verify with assertions
- More about assertions in coming lectures

```
assert property (@(posedge clk) $onehot(state_vec)) else $error("state is not one hot");
```

A Word About Memories

- Careful with memories with too many read or write ports
 - It is easy create multi ported memories:
 - “A = Mem[a], B= mem[b], C=mem[c]” implies three ports
 - Each independent address that you read/write is another port
- Tradeoff: latch rams, flop rams, register files, block memories
 - Register files: full swing (no sense-amp)
 - Better for smaller memories and generally a bit faster
 - Block memories: low swing (with sense-amp)
 - Need BIST logic / support for test mode
 - Flop/latch rams:
 - Fastest, but poor density
 - Scanned flop rams are really big

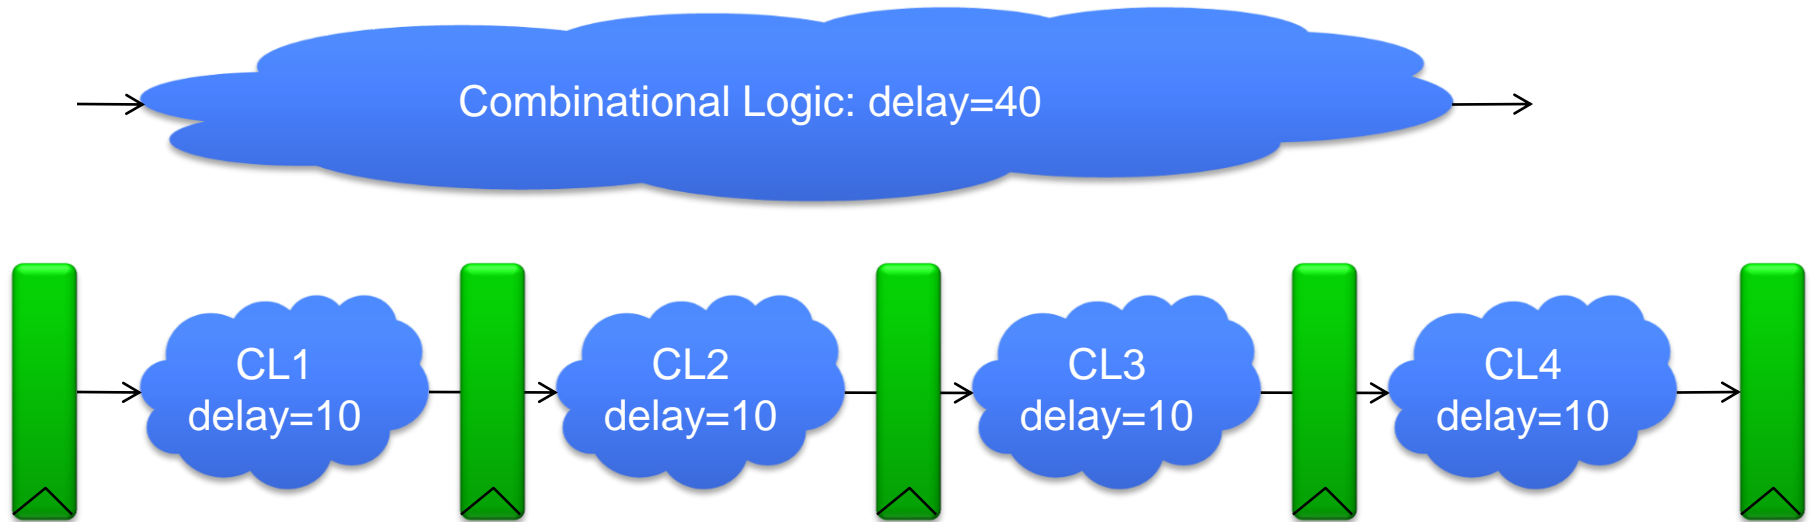
Cell Type	Area@45nm
Reg. File	1-1.5um ²
SRAM*	0.346um ²
DFF / Sdff**	4.522 / 6.118um ²
Latch**	3.458um ²

* IEDM 2007 (Intel)

** Nangate Open Cell Library

Pipelining

- Remember why pipelining is useful
 - Take a function with long delay
 - Create many functions with shorter delay that run in parallel
 - Make the latency through the function slightly worse
 - But increase the throughput of the machine (concurrency)




Limits to Pipelining


- What limits the degree of pipelining in hardware?

Retiming – A Method For Data Path Optimization

- Register location or pipe subdivision can be tricky
 - Example: where should you place an extra (middle) register?
 - Left → paths not even; middle → 8x registers; right → not even + 4x



- Solution: Don't worry about where the pipeline registers are
 - Let the tools do retiming for you
 - Tools see one big comb. cloud → can optimize cross sections



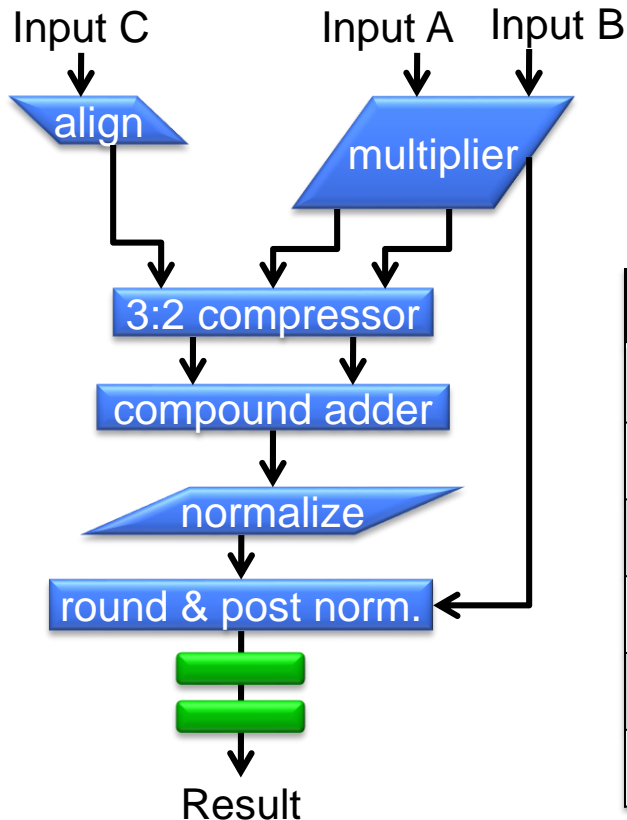
Retiming, Cont'd

- Goals
 - Clock period (min-period retiming)
 - Number of registers (min-area retiming)
 - number of registers for a target clock period (constrained min-area retiming)
- Problem definition: From circuit to a directed graph
 - Gate \leftrightarrow Vertex (V is the set of vertices)
 - Wire \leftrightarrow Edge (E is the set of edges)
 - $d(v)$ denotes delay of a gate, $d(v) \geq 0$
 - $w(e)$ denotes #registers on edge, $w(e) \geq 0$


Retiming Circuit Representation (Example)

Circuit

(simplified floating point multiply accumulate unit)




Representation Graph



OP	Delay
align	5
mult.	8
3:2	1
adder	5
norm	4
round	4

Basic Operation

- Registers can move from input to output of a gate (or vice versa)
- Does not affect gate functionalities
- Problem reduced to a graph optimization problem




Retiming Problem Formulation (Tools' Perspective)

- For a path $p: v_0 \rightarrow v_k$

Path delay is $d(p) = \sum_{i=0}^k d(v_i)$


#Cycles is $w(p) = \sum_{i=0}^{k-1} w(e_i)$

- To find the resulting clock cycle:
find the longest path that has no
flops: $c = \max_{(p|w(p)=0)} \{d(p)\}$
- For the FMA, $c = 22$

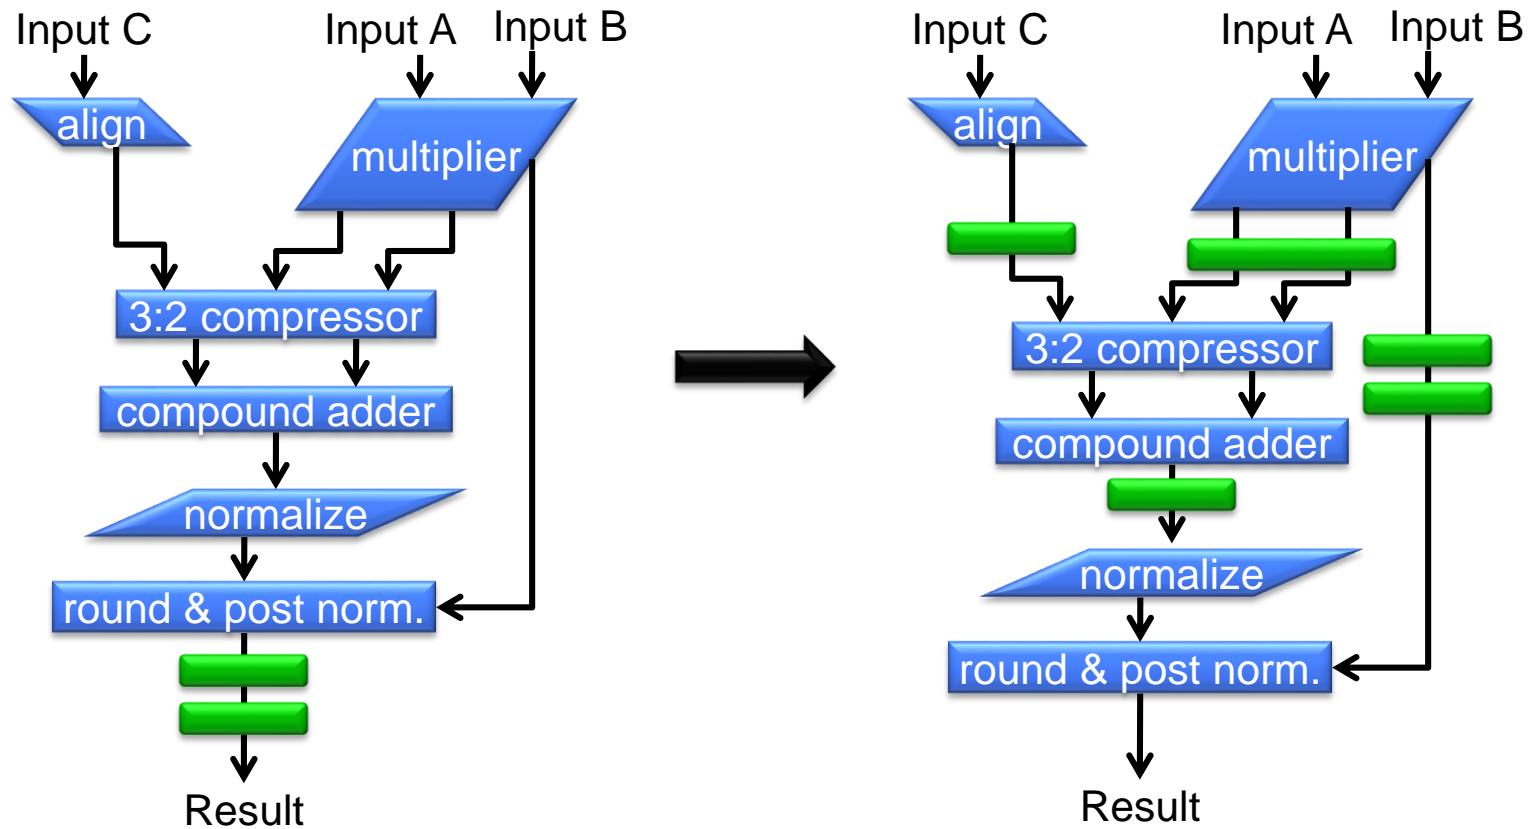


The Tool Crunch The Numbers

- Optimizing the FMA graph: $c = 8$



Retimed Netlist



Retiming In Practice Today



- Retiming considerations
 - No combinational loops. Each loop must contain ≥ 1 flops
 - Any power-up state of the design must be safely handled
- Retiming Achilles: Verification
 - Retiming does not change the signature values...
 - But it may change the cycle at which a signature appears!
 - Significantly more computationally intensive than combinational equivalence
- Verification technology needs to be improved for greater acceptance
- Mostly used in pipelined data-path

Dealing with Routing Issues Before they Become Huge Problems

- Communication can create nasty problems late in design
 - So it is good to think about them early on
 - And avoid doing things you will regret later
- Delay and power cost of long wires
- Dealing with large fan-in circuits
- Dealing with large fan-out circuits

Long Wires


- Wires introduce noise coupling (cross-talk)
 - Why?
- Wires add both resistivity and capacitance
 - Proportional to length: $R_w = \rho / (t * W) * L$ and $C_w = \epsilon * W / t * L$



- For **long wires** delay is proportional to L^2 since $D \sim R_w * C_w / 2$
 - Wires load and slow down the gate that drives them
- Adding distributed inverters to a wire can make it faster

Repeaters

- Repeaters: Break the wire into two or more pieces
 - Synthesis tools (should) do buffer insertion




- What's the optimal number of repeaters for a line?
 - Add N identical repeaters → Delay = N * Repeater delay
 - Optimize repeater's size and find the optimal N:

- Turns out it is **linear with L**:
$$N_{opt} = L \cdot \sqrt{\frac{C_w \cdot R_w}{2 \cdot C_G \cdot R_T}}$$

- This means **Delay is linear with L!**

Dealing With Large Fanin

- Hot Spots
 - Huge demand for routing resources in a very small physical space
 - RTL Structures have high # input pins per unit area
- Designer should know where such structures exist (and avoid)




Example: Large Muxes


RTL

```
reg [2:0] SEL;  
reg [127:0]  
A,B,C,D,E,F,G,H,M;  
always @ (A or B or C or D  
or E or F or G or H or  
SEL)  
begin  
  case (SEL)  
    3'b000 : M = A;  
    3'b001 : M = B;  
    3'b010 : M = C;  
    3'b011 : M = D;  
    3'b100 : M = E;  
    3'b101 : M = F;  
    3'b110 : M = G;  
    3'b111 : M = H;  
  end  
end
```

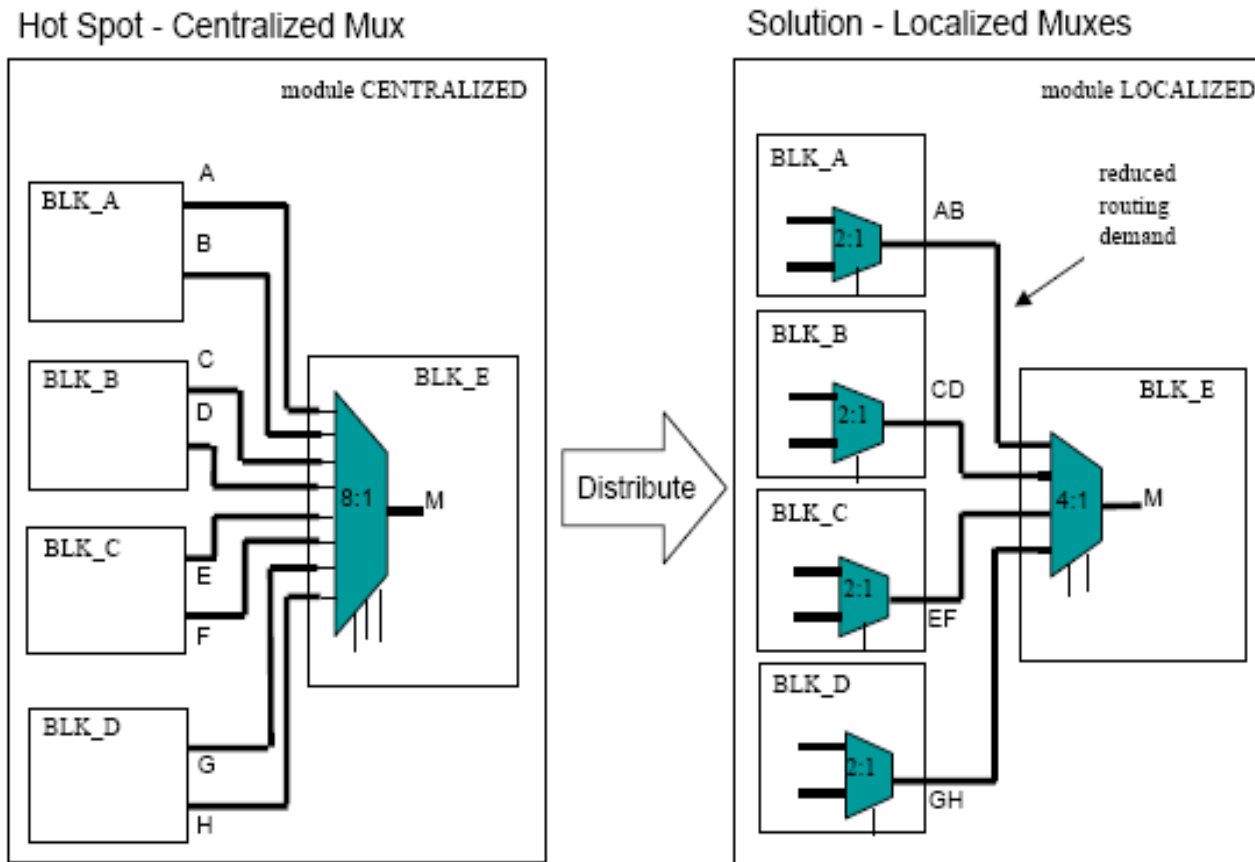
Gate-Level Netlist



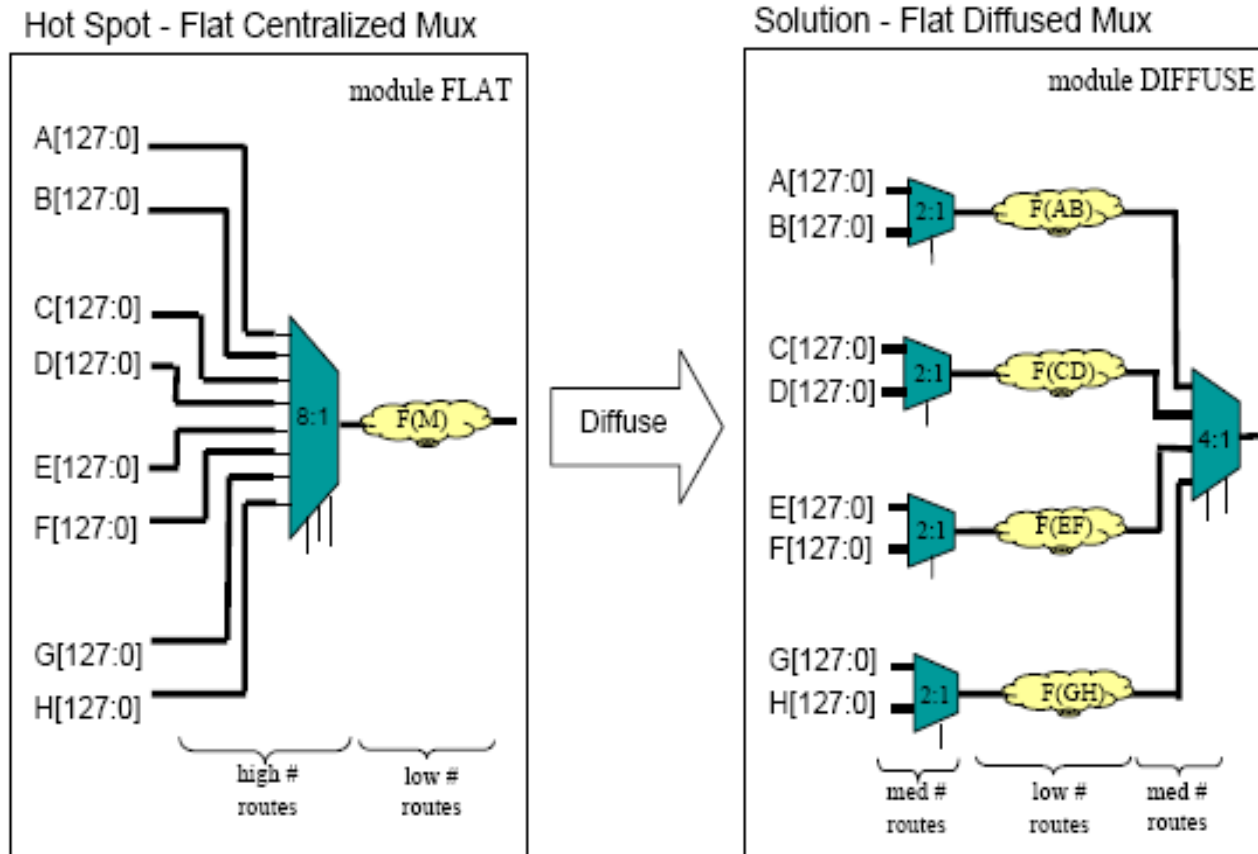
Physical Design



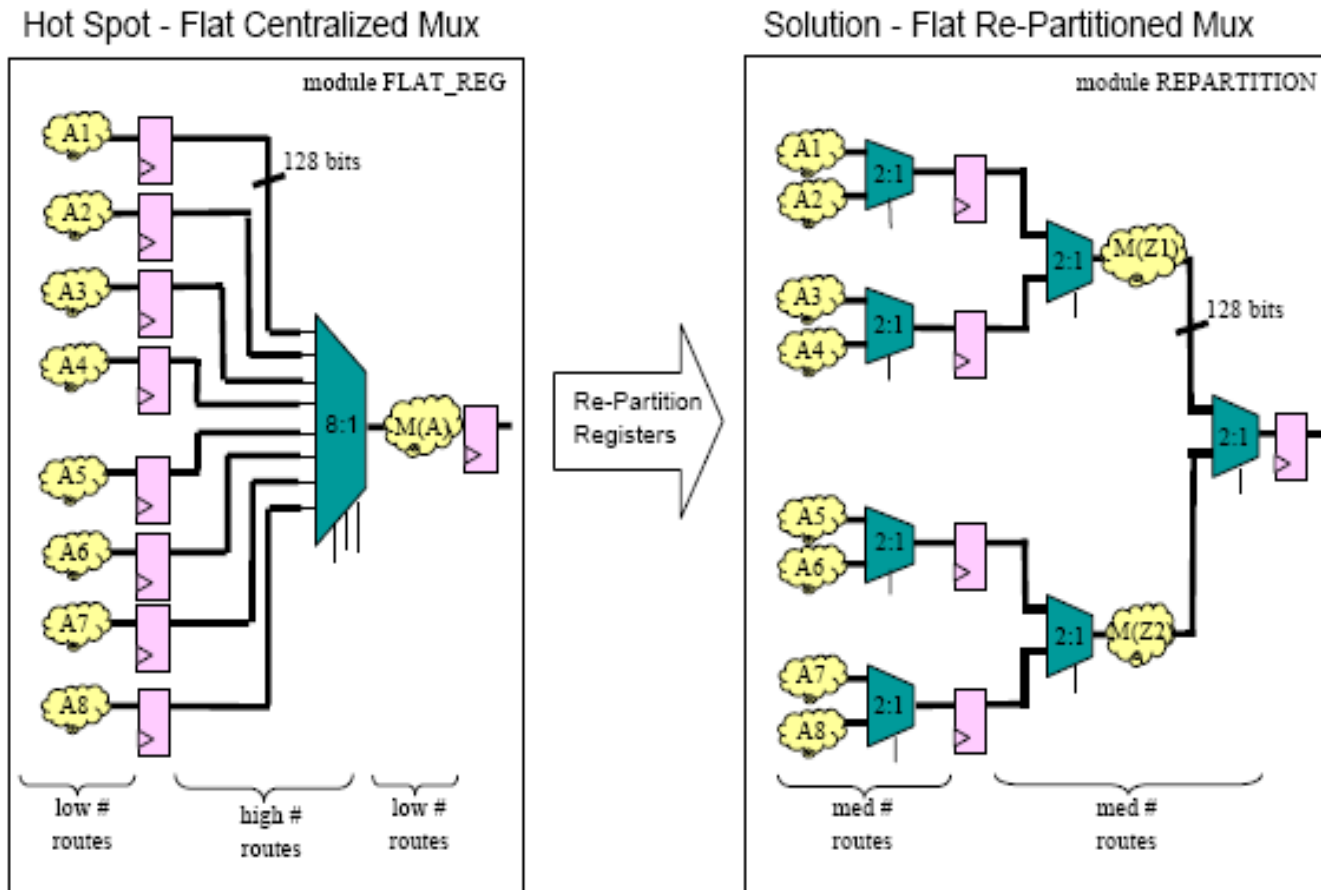
Solution 1: Hierarchical Separation



Solution 2: HW Repetition & Logic Diffusion




Solution 3: HW Repetition & Logic Diffusion & Retiming



http://www.synopsys.com/news/pubs/sjsnug/sj03/marshall_pres.pdf

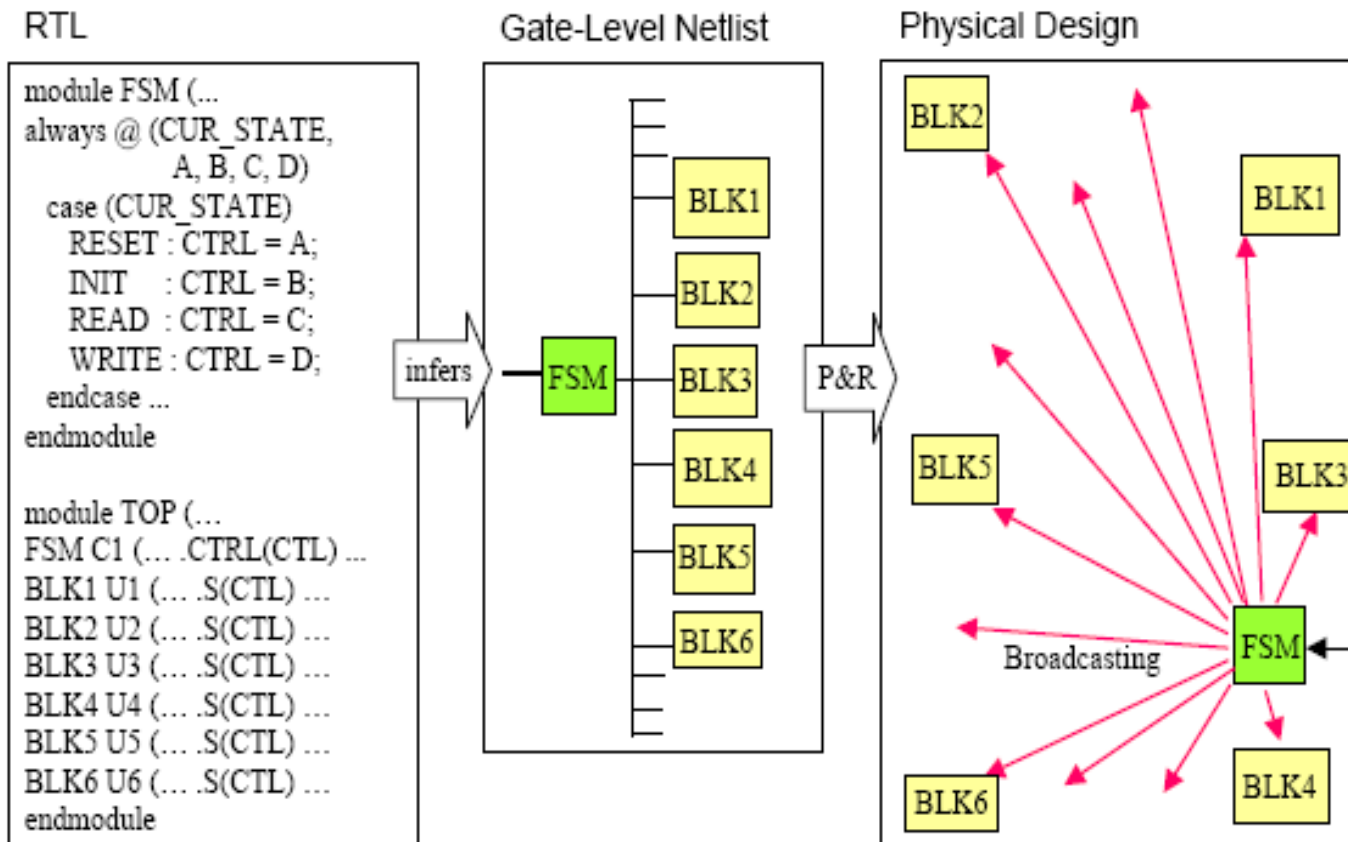
Broadcasts



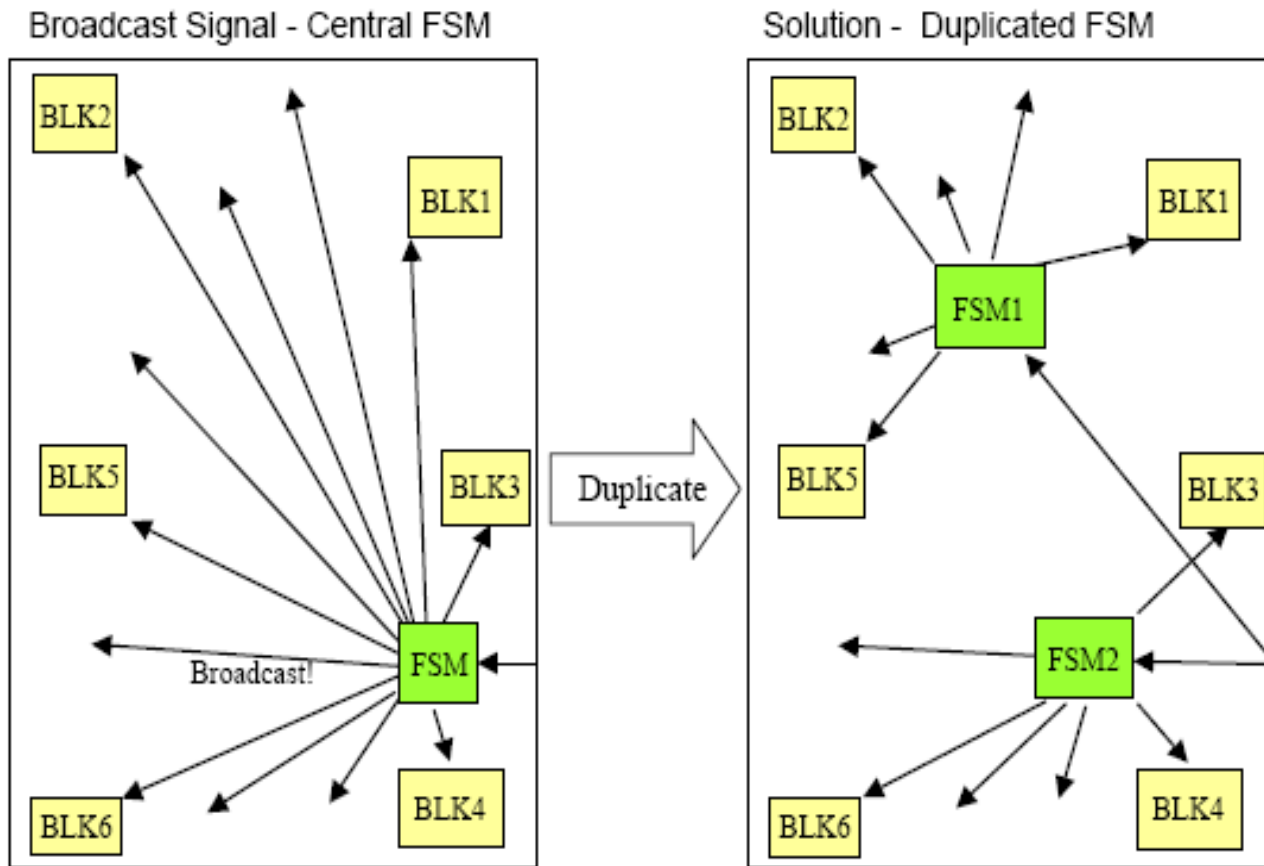
- Non clock or reset signal that has high fanout, crosses many hierarchical boundaries, and is time critical
- Long routes cause timing problem & impact routability
- Early RTL floorplanning & RTL analysis can identify signals

http://www.synopsys.com/news/pubs/sjsnug/sj03/marshall_pres.pdf

Example: Control logic



Solution: FSM Duplication



Bottom Line

- Gates are cheap
- Wires are expensive
 - Delay, energy and area, the big three
- Need to plan / optimize wires
 - It is ok to duplicate gates in most cases
 - Will end up being faster, lower power
 - And sometimes even smaller

Design For Verification

- Unit Interfaces – Be sure to make them simple and well defined
 - Litmus test: Can the module be replaced with a black box?
- Wrap any “magic” construct with appropriate compile flags
 - Make sure “magic” constructs are well documented


```
`ifdef MAGIC_MEM_INIT  
    for (i=0; i<128; i++) mem[i] = 0;  
`endif /* MAGIC_MEM_INIT */
```

- Where possible – parameterize
 - It is often useful to simulate smaller memories/queues/fifos
 - Can hit more corner cases in less simulation time
 - Enables mapping of a “reduced version” to emulation platforms

Points Where Verification and Design Interact: Assertions

- Don't wait to test your assumptions – **Add Assertions!**
 - For inputs:
 - E.g., Does your module receive requests **only** after *busy==0*?

```
Assumption1: assert property (  
    @(posedge clk) req |-> $past(busy,1)==1'b0  
)else  
    $error("ERROR – req after busy asserted");
```



- For things you're just not quite sure of
 - Remember the “one-hot” example from a few slides back...
 - For your block's output: to force protocol obedience
- Assertions are great for both uncovering bugs as well as for debugging/pinpointing them

Help Verification Sync With Design Intent

- Keep key signals/states explicit
 - In cases, the verification must know the current state

```
if (DUT.proc.init_done)
```

```
    Run_test()
```

Vs.


```
if (DUT.proc.no_interrupt & DUT.proc.alu_rdy & DUT.proc.reg_file_init & ...)
```

```
    Run_test()
```

- Similarly, use explicit arbitration signals
- Inline Comments – No one can read your mind
 - Especially if you optimized the heck out of the algorithm ;-)
 - Bonus: You'll find many bugs as you write your thoughts


Design For Post-Silicon Testing Debug Pathways - “Is it alive?”

- For key signals, and keep it as simple as possible
 - Key signals: various clocks, reset, select, etc.
 - Simple: control the debug path with minimal logic
 - Scope: unit or system level signals
- No logic should ever depend on a debug path
 - This is good news: Don't worry about the debug path timing
 - If needed, run the machine slower
- I like the following scheme:
 - Zero synchronous logic
 - Minimal comb. logic
 - Can be hierarchical mux




Internal State Visibility - “Does it work?”

- Is the memory being written? Does the configuration register hold the right value? Is the counter being updated?
- Enable external reads of the explicit state of the machine
 - Explicit state = memories, register files, machine configuration...
 - Enabling writes (in addition to reads) is useful
- Typically done through the JTAG interface
 - JTAG = Joint Test Action Group, IEEE 1149.1 standard
 - Minimal additional I/O: {TCK, TDI, TDO, TMS, TRST}

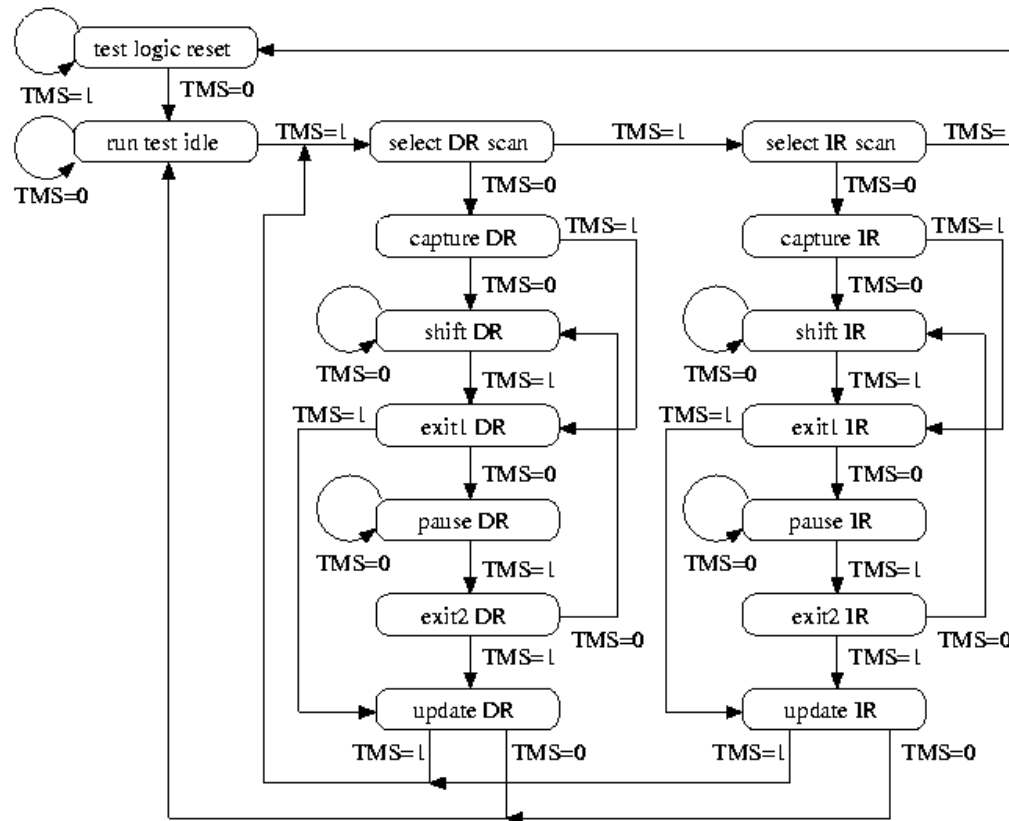


JTAG

- Works for any data width - serial interface
- Typically runs at 1/2 to 1/10 of core frequency
 - Not in critical path
 - beware of time domain crossing
- Multiple TAPs can be chained together (serially)



JTAG State Machine




JTAG Test Access Port (TAP) controller state transition diagram

Reading All Flops - “When all else fails...”

- Machine state is not only memories and register files
 - It is every flop along every path or pipeline
 - But not all of them have explicit read ports


- Conceptual view of a scan register:



- Automatic tools stitch all flops into a long chain (1000s/chain)
- The exact machine state can be read at any point of time**

Controlling All I/Os – Boundary Scan Registers

- Conceptual view of an I/O pad:



- All boundary scan registers are chained
- Controlled via JTAG:
 - SAMPLE: sample all inputs and outputs; values are then scanned out.
 - INTTEST: Force all inputs to a previously-scanned-in pattern
 - EXTEST: Force all outputs to a previously-scanned-in pattern
- What is it good for?

The End
