

## 4 Lecture: The Process Model

### Outline:

- Announcements
- From last time
- The Process (Users' View)
- The Operating System's view: The context
- Example of a context switch
- System Calls Again
- System Call Mechanisms
  - How to do it
  - And onwards
- Onwards
- The Process Model: a little deeper
- Pseudoparallelism and nondeterminism
- Possible process states

### 4.1 Announcements

- Coming attractions:

Event	Subject	Due Date			Notes
asgn2	LWP	Mon	Jan 26	23:59	
asgn3	dine	Wed	Feb 4	23:59	
lab03	problem set	Mon	Feb 9	23:59	
midterm	stuff	Wed	Feb 11		
lab04	scavenger hunt II	Wed	Feb 18	23:59	
asgn4	/dev/secret	Wed	Feb 25	23:59	
lab05	problem set	Mon	Mar 9	23:59	
asgn5	minget and minls	Wed	Mar 11	23:59	
asgn6	Yes, really	Fri	Mar 13	23:59	
final (sec01)		Fri	Mar 20	10:10	
final (sec03)		Fri	Mar 20	13:10	

Use your own discretion with respect to timing/due dates.

- Asgn1 due Wednesday (reminder about late days)
- All assignments out
- Lab02: Why minix 3.1.8

### 4.2 From last time

- History
- Make(1) in AM
- System call mechanisms
- Booting
- Regaining control

### 4.3 The Process (Users' View)

A quick overview of the (human) users' view of the system:

Each running program becomes a process, isolated from all other processes:

- Each process has the illusion of being alone
- Has its own memory
- Has its own scheduling time
- Has its own interface to the outside world (file descriptors)
- All interprocess (and outside world) interaction takes place through the operating system.

Processes have identity:

- User ID
- Group ID
- Process ID

Processes have resources:

- memory (address space)
- time

### 4.4 The Operating System's view: The context

In the OS's view, a process consists of some resources:

<b>registers</b>	Each process gets its own copy.
<b>address space</b>	A region of memory, usually starting at address 0, corresponding to a particular process.
<b>identity</b>	uid, gid, pid, ppid, etc. All those things that determine a process's identity and privilege.
<b>file descriptors</b>	Connections to the global file descriptor table to allow for IO
<b>signals/masks</b>	Pending notifications

The OS's purpose is to keep these separate. How?

- Records kept in a **process table**
- Processes run for an allotted time(**quantum**), or until they yield (e.g. disk wait)
- **context switch**
  - Old process suspended (timer?)
  - Old process's registers saved (where?)
  - Old process's memory—text, stack, data—saved in a **core image**
  - Memory set up for new process (text, stack, data)
  - Registers set up for new process
  - New process “continued” as if nothing had happened.  
Imaging blinking and discovering that it was three hours later.

## 4.5 Example of a context switch

How does it really happen? (with pictures, and everything.)

- Process A is running (Figure 4.5a)
- An interrupt occurs (Figure 4.5b)
- Push registers to preserve them. (Figure 4.5c)
- Save SP in the process table and switch to the operating system (kernel) stack (Figure 4.5d)
- The OS is now running:
  - Preserve Process A's memory (if necessary)
  - Choose the next process to run.
  - Load Process B's memory (if necessary)
- Restore B's SP from process table. (Figure 4.5e)
- Pop B's registers (Figure 4.5f)
- Return from interrupt (Figure 4.5g)

## 4.6 System Calls Again

We said “An OS is defined by its system calls”. What does that mean?

### System Call

A system call is the means by which the kernel provides access to a particular operating system service, and the services available through system calls (e.g., reading and writing the disk, allocating memory, starting new processes, etc.) are reserved to the kernel; there is no other way of doing these things.

## 4.7 System Call Mechanisms

Last time we said

- It's all about privilege.

### 4.7.1 How to do it

How done? Machine dependent. usually in assembly, but hidden from the users' view.

In any case, some sort of trap is involved to get supervisor privileges.

A *trap* is a software generated interrupt. Exactly what happens varies from architecture to architecture, but the result usually involves:

- saving the state of the currently executing program
- elevating of processor privilege level to “supervisor”



Figure 5: The process of a context switch  
35

- transfer of control to a pre-registered routine
- *The OS does something*
- Then, usually:
  - privilege level is reduced back to user
  - original state of executing program is restored

#### 4.7.2 And onwards

- How does the OS get control back?

### 4.8 Onwards

Last time we talked about context switching and the concept of a Process. Now that we have the general idea, let's look at some general structures.

This week we are going to look at general OS architectures, then move into scheduling.

### 4.9 The Process Model: a little deeper

Now, we need to talk about some details of how it's actually done and the issues involved.

The **process** is the most important concept in understanding operating systems.

The operating system provides **multiprogramming**:

**pseudoparallelism** time-sliced parallelism on a uniprocessor

(yield?, pre-empt?)

**true parallelism** provided by a multiprocessor

Keeping track of parallel activities is difficult, so OS designers have developed the model of the **sequential process**:

Each process (**running program**) has its own:

- IO Access: open file descriptors
- Register file
- Memory
  - Text (code) segment
  - Data (bss and heap) segment
  - Stack segment
- Program Counter

### 4.10 Pseudoparallelism and nondeterminism

(Ir?)regular (random) context switches mean programs can make no assumptions about:

- order of execution relative to other programs
- timing (example: a clock, or game timing)

(If timing *is* important, you need to get a Real-Time O.S.)

## 4.11 Possible process states

Not all processes can run at the same time:

At any given time, any process can be in one of three possible states: Running, Ready, or Blocked. Possible transitions between these states are shown in Figure 6.




Figure 6: Possible states for a process

In more detail, a process can be:

**Running** The process is loaded in memory and is currently executing on the cpu.

**Ready** The process is runnable, but another process has the cpu.

**Blocked** The process is waiting for some external event to enable it to run. E.g. completion of some IO event, or an alarm, or availability of more memory.