

# EE 531: ADVANCED VLSI DESIGN

---

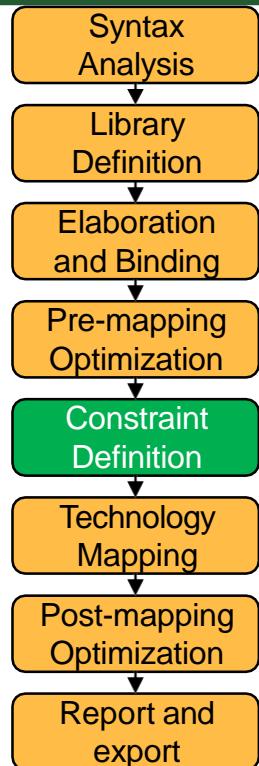
## Logic Synthesis: Part 2

Nishith N. Chakraborty

January, 2025

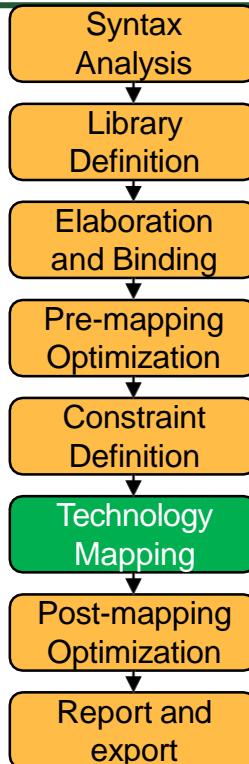
# CONSTRAINT DEFINITION

- Following **Elaboration**, the design is loaded into the synthesis tool and stored inside a data structure.
- Hierarchical ports (inputs/outputs) and registers can be accessed by name.
- At this point, we can load the design constraints in SDC format.
- Carefully check that all constraints were accepted by the tool!



# TECHNOLOGY MAPPING

- Technology mapping is the phase of logic synthesis when gates are selected from a technology library to implement the circuit.
- Why technology mapping?
  - Straight implementation may not be good.
  - For example,  $F=abcdef$  as a 6-input AND gate causes a long delay.
  - Gates in the library are pre-designed, they are usually optimized in terms of area, delay, power, etc.
  - Fastest gates along the critical path, area-efficient gates (combination) off the critical path.
- Can apply a minimum cost tree-covering algorithm to solve this problem.



# TECHNOLOGY MAPPING ALGORITHM

---

- Using a **recursive tree-covering algorithm**, we can easily, and almost optimally, **map a logic network** to a **technology library**.
- This process incurs three steps:
  - **Map netlist** and tech library to simple gates
    - Describe the netlist with only NAND2 and NOT gates
    - Describe SC library with NAND2 and NOT gates and associate a **cost** with each gate
  - **Tree-ifyng** the input netlist
    - Tree covering can only be applied to trees!
    - Split tree at all places, where fanout  $\geq 2$
  - **Minimum Cost Tree** matching
    - For each node in your tree, recursively find the minimum cost target pattern at that node.
- Let us briefly go through these steps

# STEP 1: SIMPLE GATE MAPPING

- Apply De Morgan laws to your Boolean function to make it a collection of NAND2 and NOT gates.
- Let's take the example of multi-level logic minimization:

$$\begin{aligned}
 t_1 &= d + e; \\
 t_2 &= b + h; \\
 t_3 &= a t_2 + c; \\
 t_4 &= t_1 t_3 + fgh; \\
 F &= t_4' ;
 \end{aligned}$$

$$t_1 = d + e = \text{NAND}(\bar{d}, \bar{e})$$

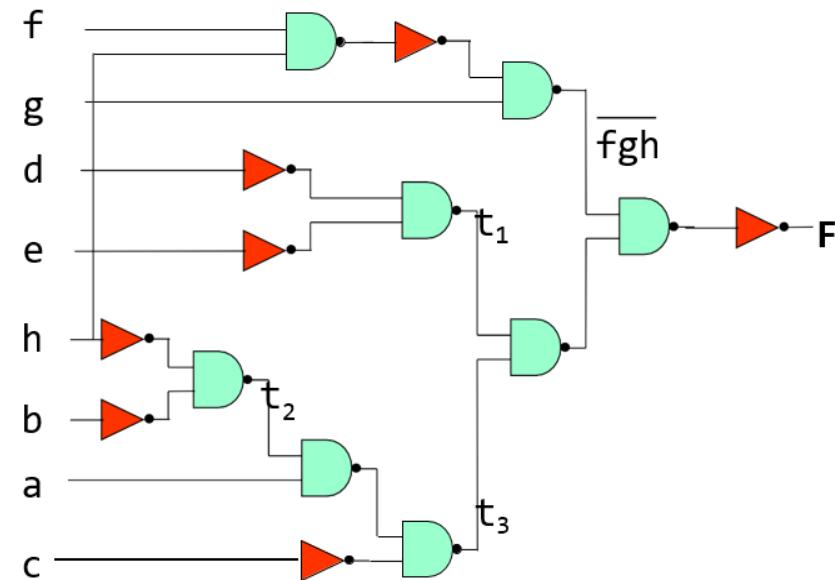
$$t_2 = b + h = \text{NAND}(\bar{b}, \bar{h})$$

$$t_3 = a t_2 + c = \overline{\overline{at_2} \cdot \bar{c}} = \text{NAND}(\text{NAND}(a, t_2), \bar{c})$$

$$t_4 = t_1 t_3 + fgh = \text{NAND}(\overline{t_1 t_3}, fgh)$$

$$fgh = \overline{\overline{fh} \cdot g} = \overline{\overline{fh} \cdot g} = \overline{\text{NAND}(\overline{\text{NAND}(f, h)}, g)}$$

$$F = \overline{t_4}$$

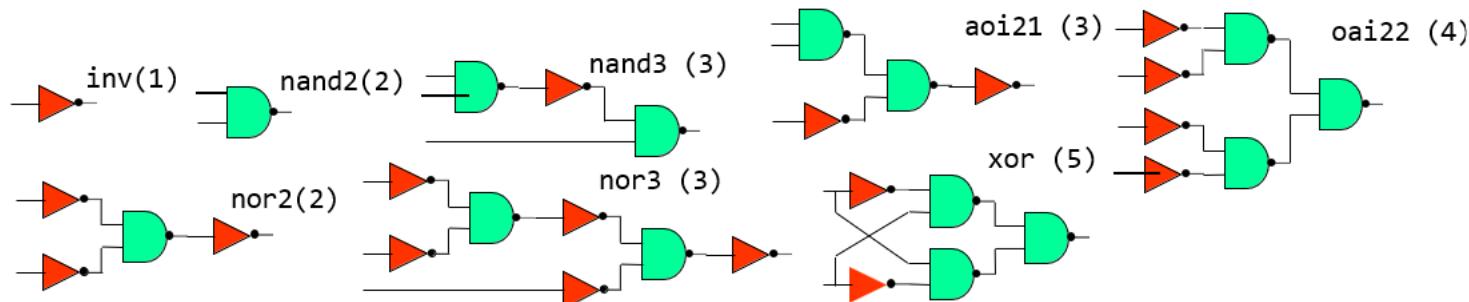


# STEP 1: SIMPLE GATE MAPPING

- And then, given a set of gates (standard cell library) with cost metrics (area/delay/power):

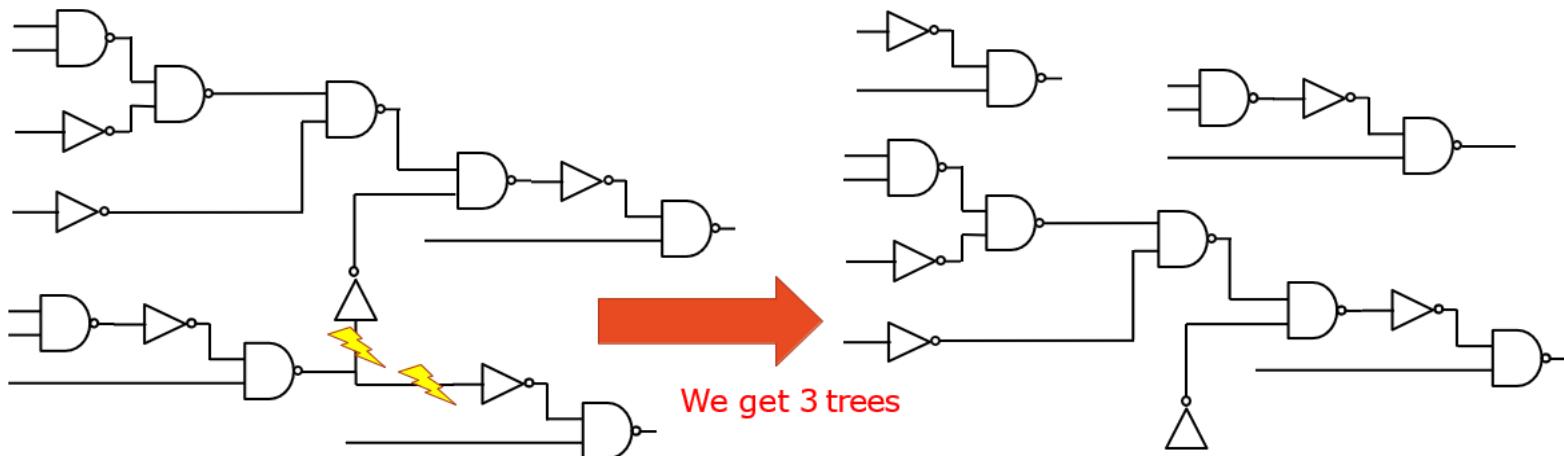


- We need to define the gates with the same NAND2/NOT set:



## STEP 2: TREE-IFYING

- To apply a tree covering algorithm, we must work on a tree!
  - Is any given logic network a tree?
  - No!
  - We must break the tree at any node with fanout $\geq 2$



## STEP 3: MINIMUM TREE COVERING

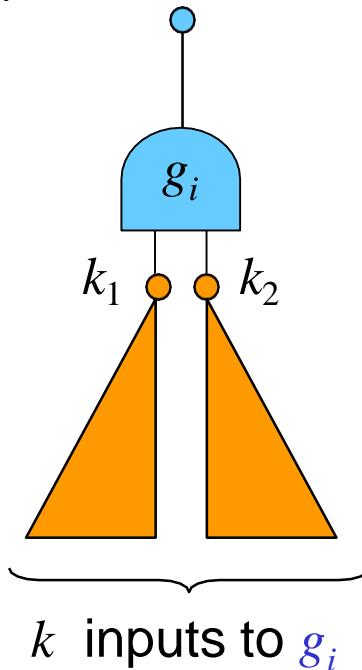
- Now, we can apply a recursive algorithm to achieve a minimum cover:

- Start at the output of the graph.
- For each node, find all the matching target patterns.

- The cost of node  $i$  for using gate  $g$  is:

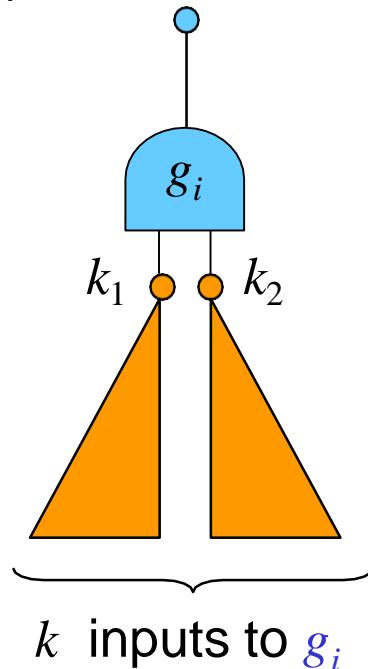
$$\text{cost}(i) = \min_k \left\{ \text{cost}(g_i) + \sum \text{cost}(k_i) \right\}$$

- where  $k_i$  are the inputs to gate  $g$ .



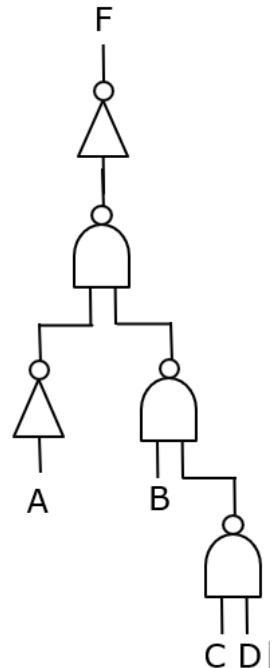
## STEP 3: MINIMUM TREE COVERING

- Now, we can apply a recursive algorithm to achieve a minimum cover:
  - Start at the output of the graph.
  - For each node, find all the matching target patterns.
  - The cost of node  $i$  for using gate  $g$  is:
$$\text{cost}(i) = \min_k \left\{ \text{cost}(g_i) + \sum \text{cost}(k_i) \right\}$$
  - where  $k_i$  are the inputs to gate  $g$ .
- For simplicity, we will redraw our graph and show an example:
  - Every NOT is just an empty circle:  I
  - Every NAND is just a full circle:  N
  - Every input is just a box:  A

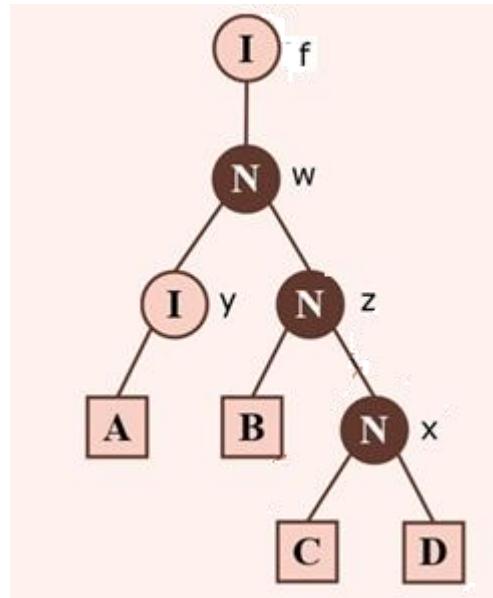
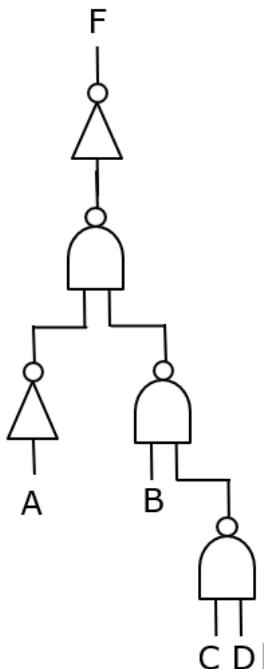


# MINIMUM TREE COVERING: EXAMPLE

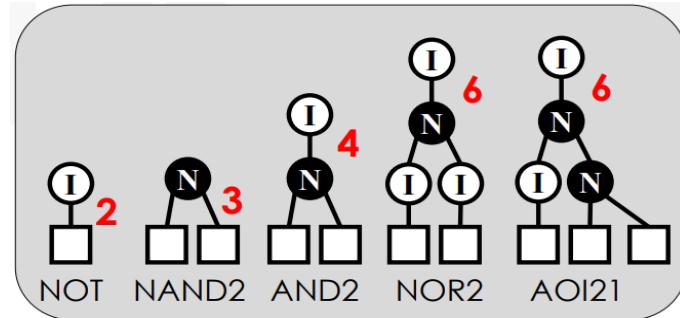
---



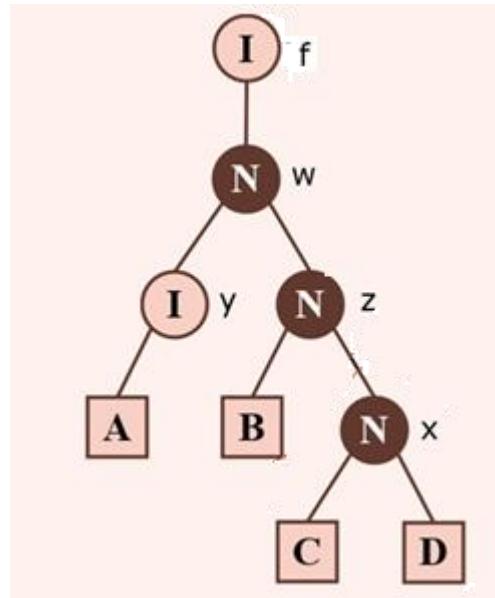
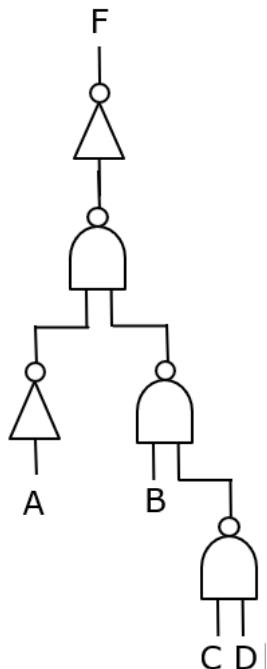
# MINIMUM TREE COVERING: EXAMPLE



**f: NOT**  $2 + \min(w)$   
**AND2**  $4 + \min(y) + \min(z)$   
**AOI21**  $6 + \min(x)$



# MINIMUM TREE COVERING: EXAMPLE



**f:** NOT  $2 + \min(w)$

**AND2**  $4 + \min(y) + \min(z)$

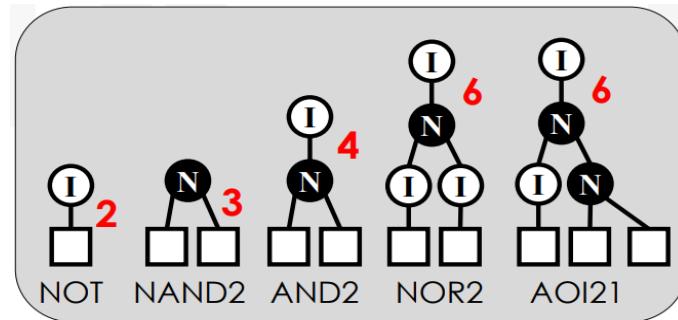
**AOI21**  $6 + \min(x)$

**w:** NAND2  $3 + \min(y) + \min(z) = 3 + 2 + 6 = 11$

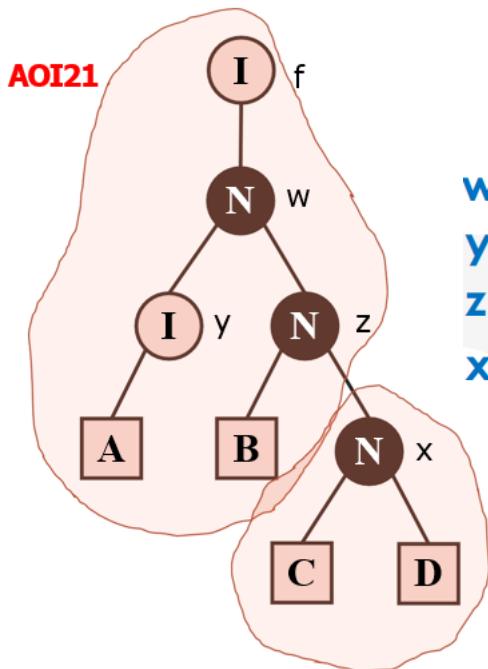
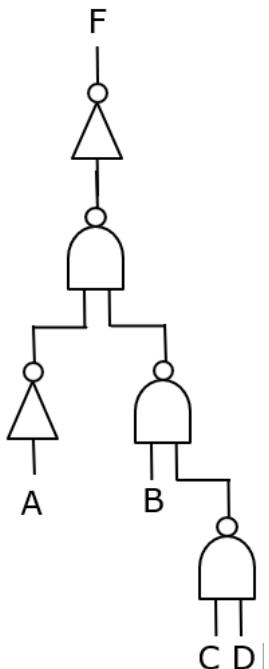
**y:** NOT 2

**z:** NAND2  $3 + \min(x) = 3 + 3 = 6$

**x:** NAND2 3

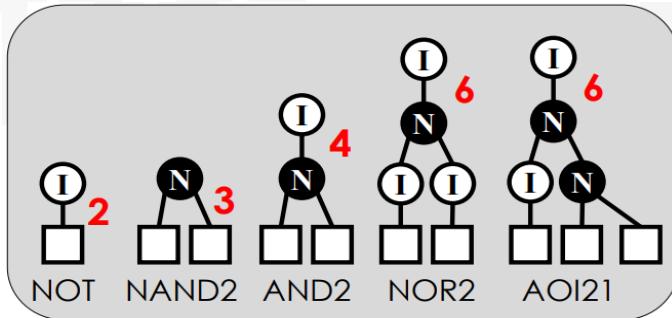


# MINIMUM TREE COVERING: EXAMPLE



**f:** NOT     $2 + \min(w) = 2 + 11 = 13$   
**AND2**     $4 + \min(y) + \min(z) = 4 + 2 + 6 = 12$   
**AOI21**     $6 + \min(x) = 6 + 3 = 9$   
**w:** NAND2     $3 + \min(y) + \min(z) = 3 + 2 + 6 = 11$   
**y:** NOT    2  
**z:** NAND2     $3 + \min(x) = 3 + 3 = 6$   
**x:** NAND2    3

NAND2

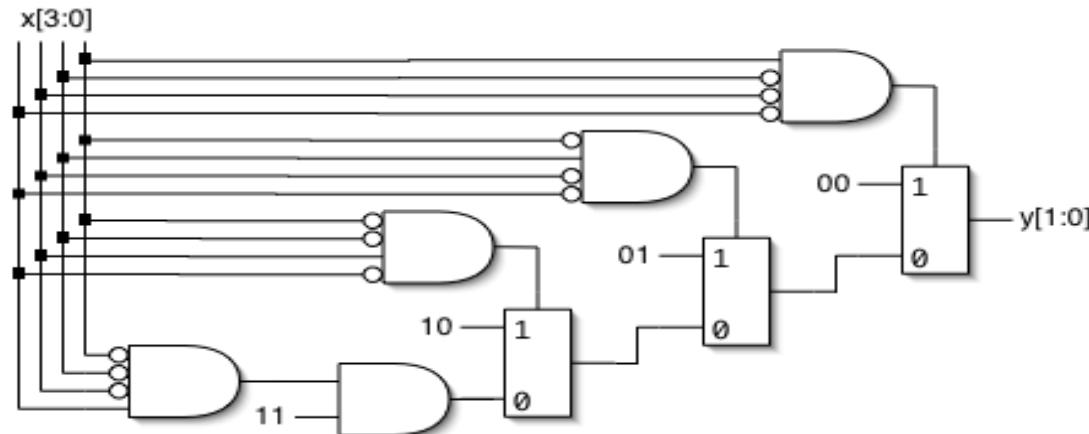


# Verilog for Synthesis: Revisited

# VERILOG TO CIRCUIT

- Let's take a simple 4→2 encoder as an example:
  - Take a one-hot encoded vector and output the position of the '1' bit.
  - One possibility would be to describe this logic with a nested if-else block:
    - The result is known as “**priority logic**”

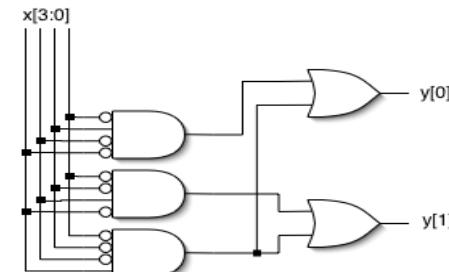
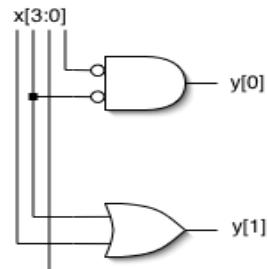
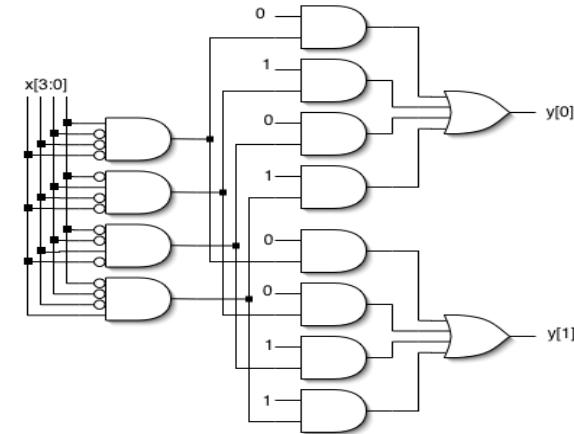
```
always @(x)
begin : encode
    if      (x == 4'b0001) y = 2'b00;
    else if (x == 4'b0010) y = 2'b01;
    else if (x == 4'b0100) y = 2'b10;
    else if (x == 4'b1000) y = 2'b11;
    else y = 2'bxx;
end
```



# VERILOG TO CIRCUIT

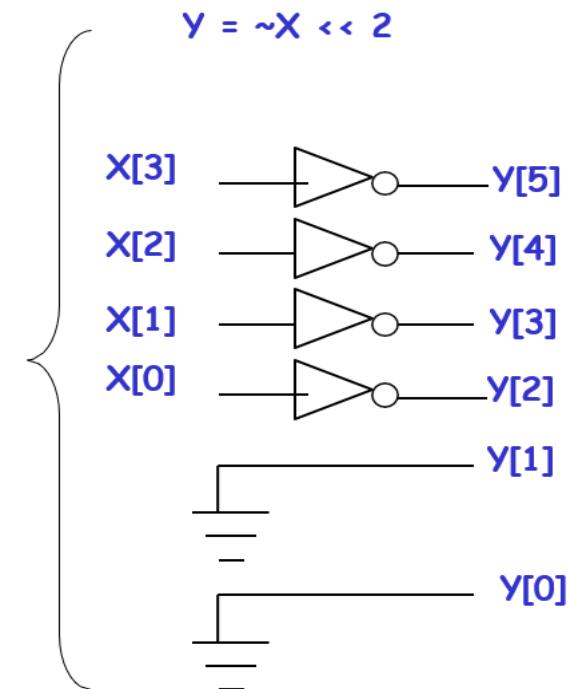
- It would have been better to use a **case** construct:
  - All cases are matched in **parallel**
  - And better yet, synthesis can **optimize** away the constants and other Boolean equalities:

```
always @(x)
begin : encode
    case (x)
        4'b0001: y = 2'b00;
        4'b0010: y = 2'b01;
        4'b0100: y = 2'b10;
        4'b1000: y = 2'b11;
    default: y = 2'bxx;
    endcase
end
```



# A FEW POINTS ABOUT OPERATORS

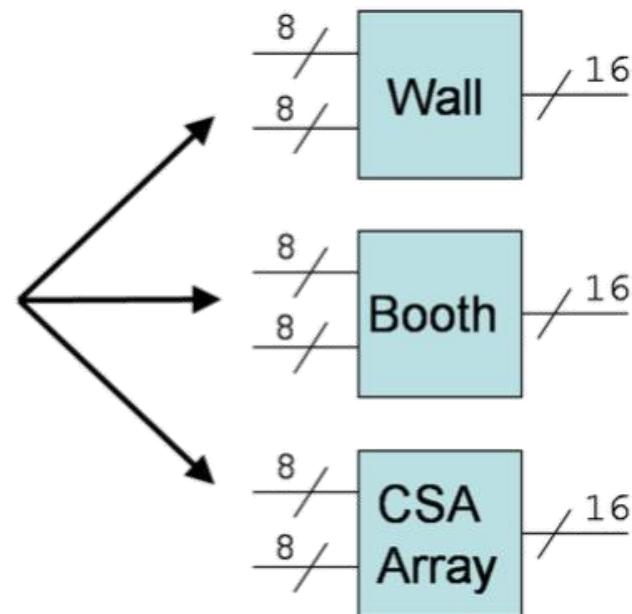
- Logical operators map into primitive logic gates
- Arithmetic operators map into adders, subtractors, ...
  - Unsigned or signed 2's complement
  - Model carry: target is one-bit wider than source
  - Watch out for \*, %, and /
- Relational operators generate comparators
- Shifts by constant amount are just wire connections
  - No logic involved
- Variable shift amounts a whole different story → shifter
- Conditional expression generates logic or MUX



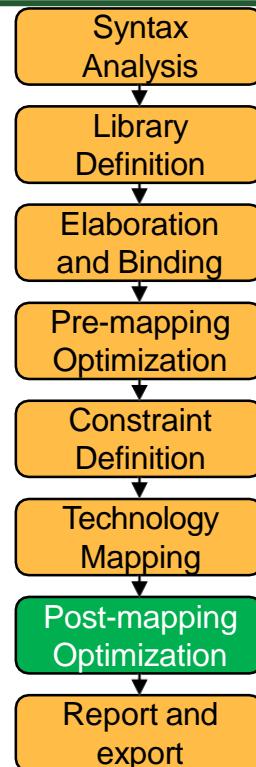
# DATAPATH SYNTHESIS

- Complex operators (Adders, Multipliers, etc.) are implemented in a special way
  - E.g. Type of multiplier

```
module mult (d1, d2, do);
  input [7:0] d1, d2;
  output [15:0] do;
  assign do = d1 * d2;
endmodule
```



# OPTIMIZATION



# OPTIMIZATION OPTIONS FOR SYNTHESIS

---

- Another look at RTL to physical layout (GDSII\*) design flows
- Metrics used for optimization
  - Speed
  - Area
  - Power
  - Others?
- ASIC/SoC design choices that allow optimizations

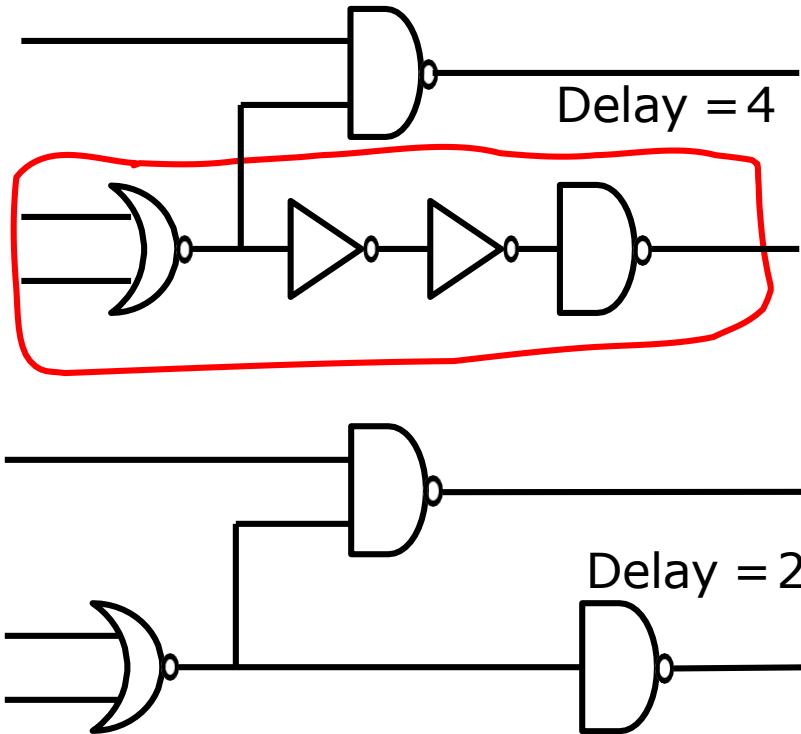
# TIMING DRIVEN SYNTHESIS

---

- Traditionally, timing has been one of the major drivers in ASIC/SoC design
- The goal is to minimize the delay on the *critical path* of the design so that the frequency is maximized
- CAD tools must include accurate and robust models for estimating the delay through the circuit(s)
- During synthesis, several choices can be made to reduce delay:
  - Optimize number of logic levels
  - Logic family used (if in standard cell library)
  - Cell sizing (usually have multiple sized std. cells)

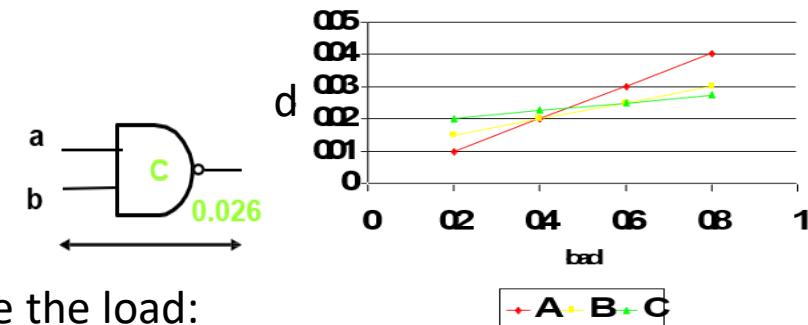
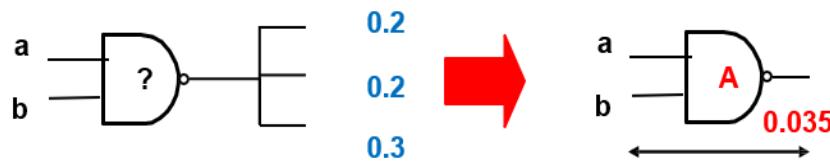
# HOW TO OPTIMIZE TIMING?

- There are many ‘transforms’ that the synthesizer applies to the logic to improve the cost function:
  - Resize cells
  - Buffer or clone to reduce load on critical nets
  - Decompose large cells
  - Swap connections on commutative pins or among equivalent nets
  - Move critical signals forward
  - Pad early paths
  - Area recovery
- Simple example:
  - Double inverter removal transform:



# RESIZING, CLONING AND BUFFERING

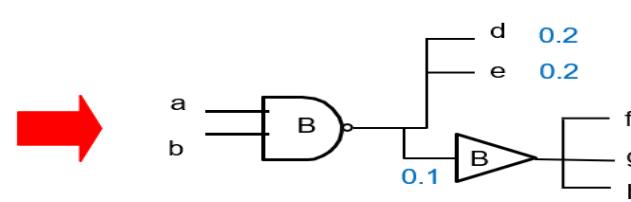
- Resize a logic gate to better drive a load:



- Or make a copy (clone of the gate) to distribute the load:

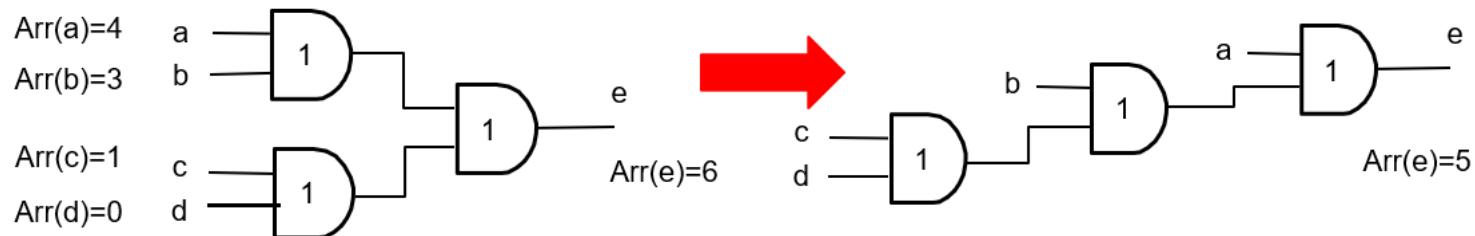


- Or just buffer the fanout net:

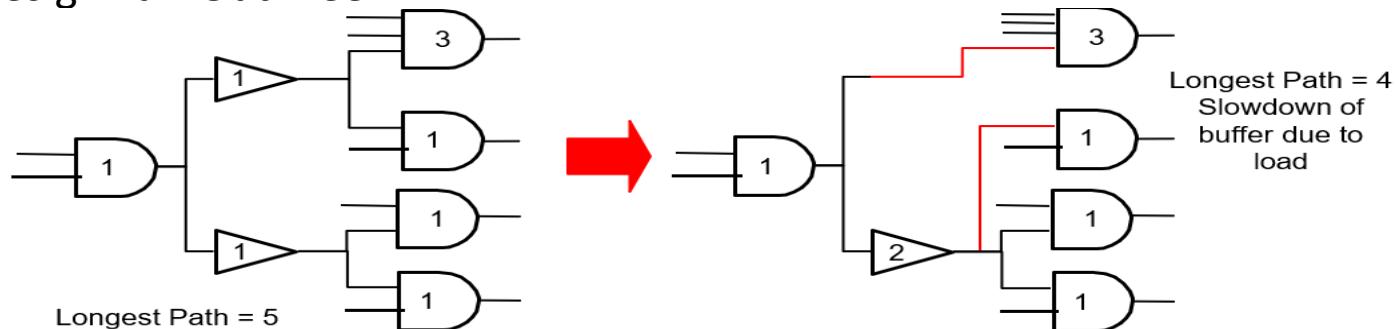


# REDESIGN FAN-IN/FAN-OUT TREES

- Redesign Fan-In Tree

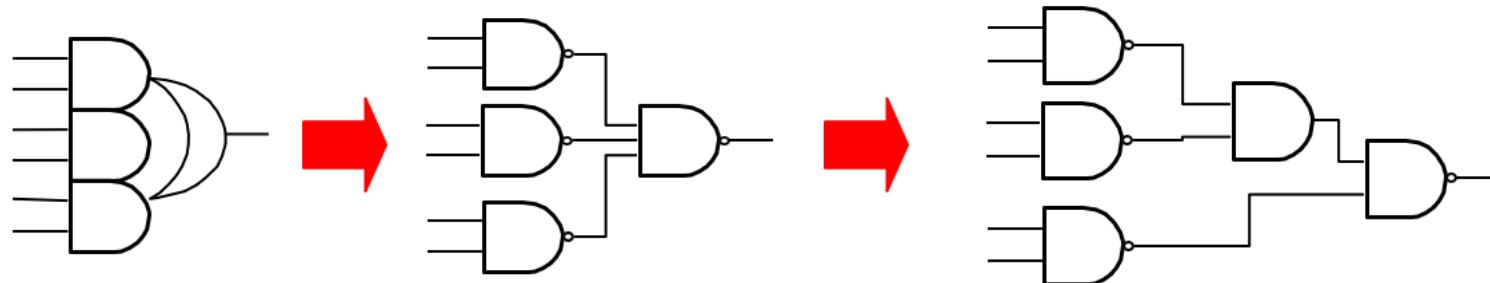


- Redesign Fan-Out Tree



# DECOMPOSITION AND SWAPPING

- Consider decomposing complex gates into less complex ones:



- Swap commutative pins:
  - Simple sorting on arrival times and delays can help



# WORST NEGATIVE SLACK

---

- Design Vision (and many other tools) uses **Worst Negative Slack (WNS)** to achieve timing requirements
- Slack = Design Delay – Predicted Delay
  - Design Frequency is essentially the target which is usually higher than Market Frequency
  - Predicted Frequency is the frequency of the current design determined by low-level simulation
- Negative slack occurs when the design does not meet the timing requirements
- Worst Negative Slack refers to the critical path, the path with the most delay

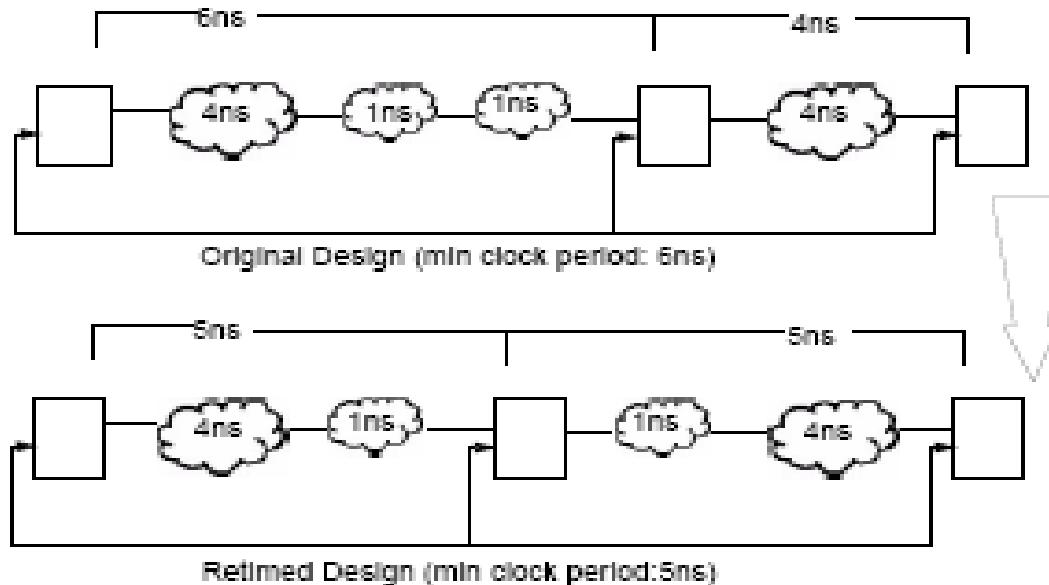
# RETIMING THE DESIGN

---

- **Retiming:** technique for improving performance of sequential circuits by repositioning registers
  - Reduces cycle time or area with no I/O latency change
- Pipelining is a subset of retiming
- Retiming redistributes sequential elements at appropriate locations to meet requirements
- Retiming does not change combinational logic

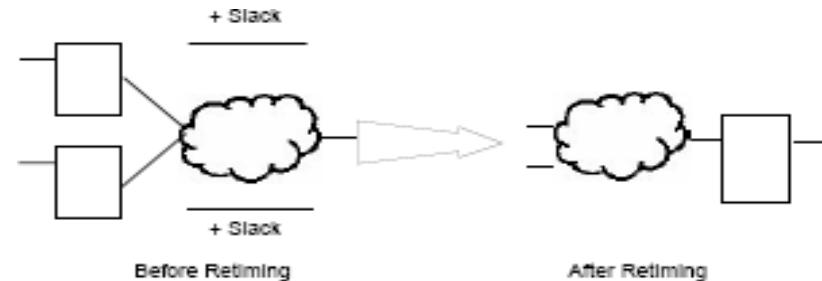
# RETIMING FOR TIMING

- Improving clock period or timing slack common
- Design Compiler distributes the registers within the design to provide minimum cycle time



# RETIMING FOR AREA

- When retiming for area, Design Vision moves registers to minimize register count without worsening the critical path in the design



# POWER AS A DESIGN METRIC

---

- Power determined by four major factors:

- Capacitance being driven ( $C$ )
- Voltage ( $V_{DD}$ )
- Frequency ( $f$ )
- Activity factor ( $\alpha$ )

$$P = \alpha \cdot C \cdot V_{DD}^2 \cdot f$$

- Low-power design techniques focus on these factors for controlling the power consumption of a design

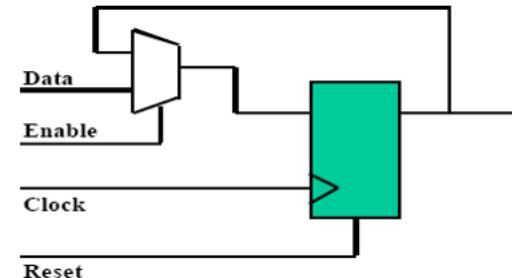
# LOW-POWER DESIGN TECHNIQUES

---

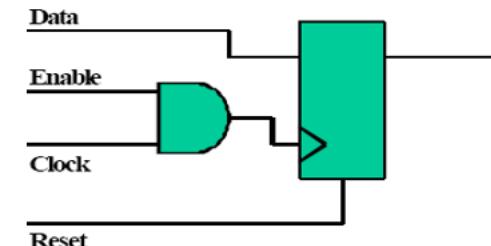
- Dynamic Voltage Scaling (DVS)
  - lower  $V_{DD}$  during runtime for quadratic savings
- Frequency scaling
- Sleep mode transistors
- Note: These are all techniques for reducing dynamic power; as technology scales, static power is becoming more of a concern

# CLOCK GATING

- Dynamic power control through synthesis typically due to clock gating
- Usually this means shutting off the clock to flip flop(s)
- Example to the right:
  - Conceptually the same
  - Implementation 1 clocks the flip flop every cycle
  - Implementation 2 only clocks when enabled
    - the lower power design



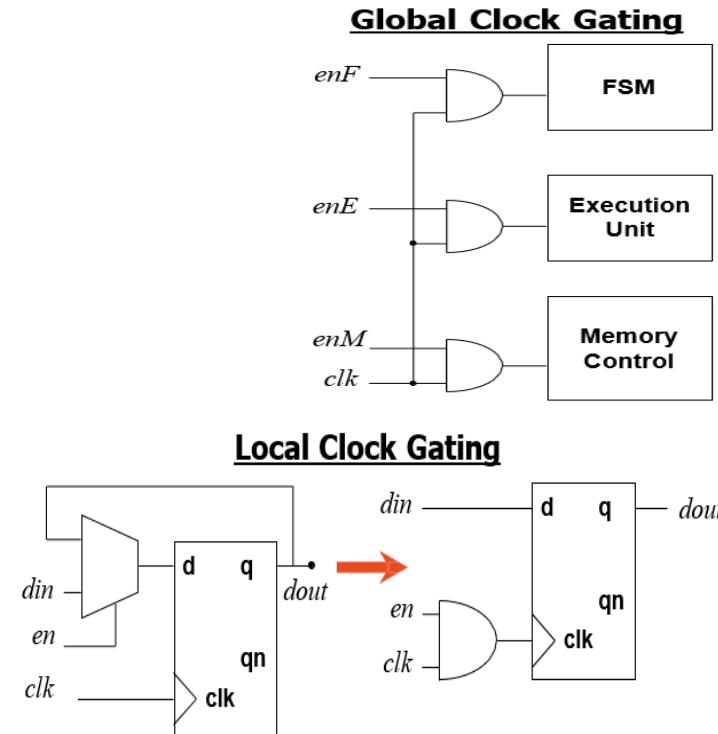
**Implementation 1**



**Implementation 2**

# CLOCK GATING LEVELS

- Block level (Global) clock-gating
  - If certain operating modes do not use an entire module/component, a clock gate should be defined in the RTL.
- Register level (Local) clock-gating
  - However, even at the register level, if a flip-flop doesn't change its output, internal power is still dissipated due to the clock toggling.
  - This is very typical of an enabled signal sampling and therefore can be automatically detected and gated by the synthesis tool.



# SYNTHESIZED CLOCK GATING

---

- To implement clock gating for power control during synthesis, tools analyze design at elaboration stage
- Most likely, gating structures are not applied to every register – cost in power of gating would exceed savings on the flip flops
- Synthesis tool tries to find gating enable signals within the design that can control the clock for a register bank

# CLOCK GATING INSTANTIATION

- Local clock gating: 3 methods
  - Logic synthesizer finds and implements local gating opportunities
  - RTL code explicitly specifies clock gating
  - Clock gating cell explicitly instantiated in RTL
- Global clock gating: 2 methods
  - RTL code explicitly specifies clock gating
  - Clock gating cell explicitly instantiated in RTL

- **Conventional RTL Code**

```
//always clock the register
always @ (posedge clk) begin
    if (enable) q <= din;
end
```

- **Low Power Clock Gated RTL**

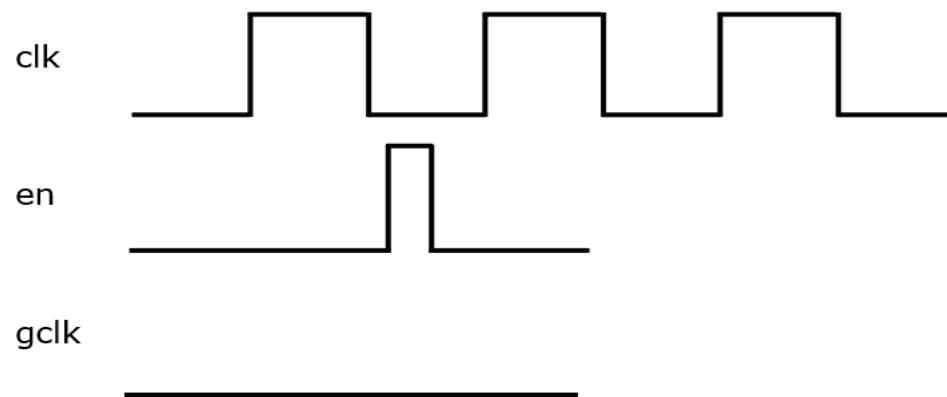
```
//only clock the ff when enable is true
assign gclk = enable && clk;
always @ (posedge gclk) begin
    q <= din;
end
```

- **Instantiated Clock Gating Cell**

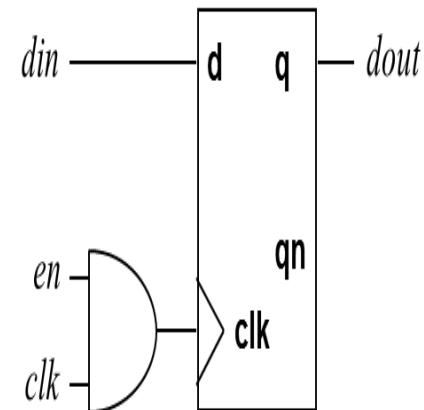
```
//instantiate a clock gating cell
clkgen#(.en(enable), .cp(clk), .gclk_out(gclk))
always @ (posedge gclk) begin
    q <= din;
end
```

# CLOCK GATING – GLITCH PROBLEM

- What happens if there is a glitch on the enable signal?

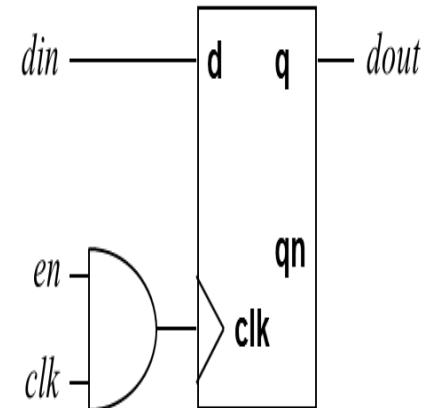
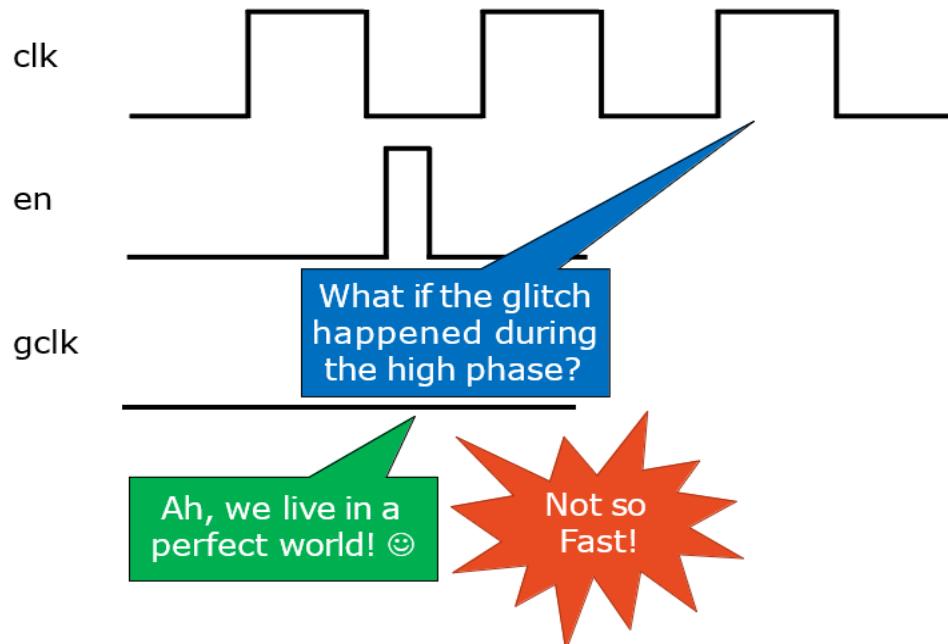


Ah, we live in a perfect world! ☺



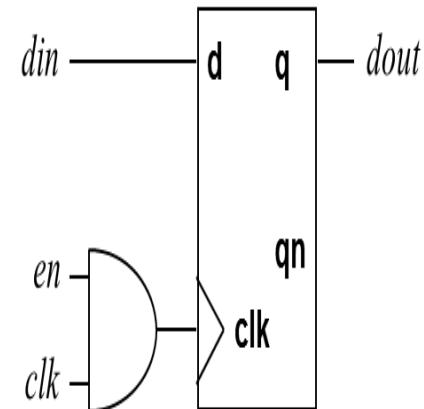
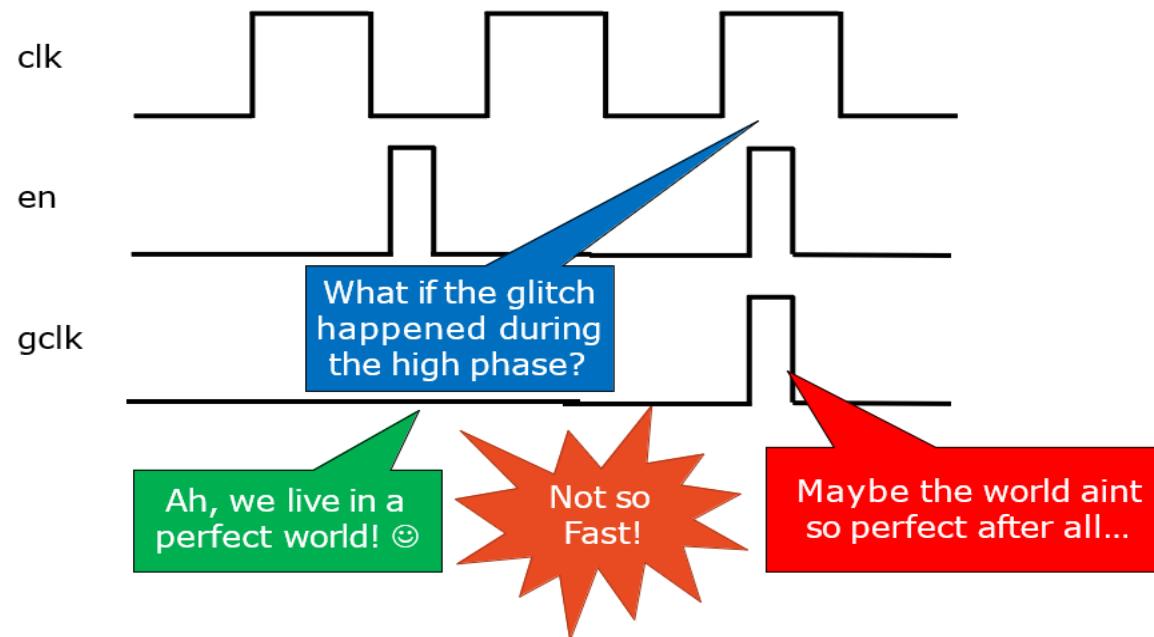
# CLOCK GATING – GLITCH PROBLEM

- What happens if there is a glitch on the enable signal?



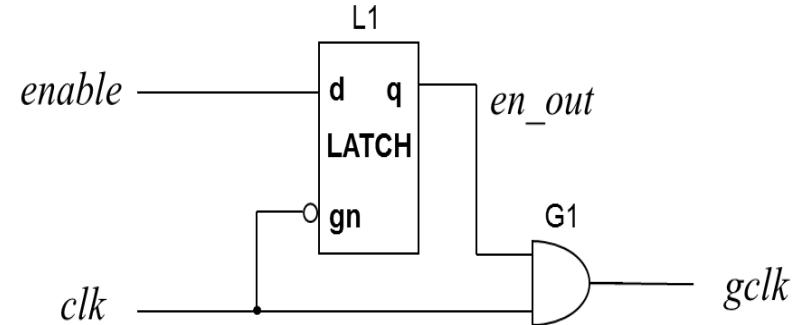
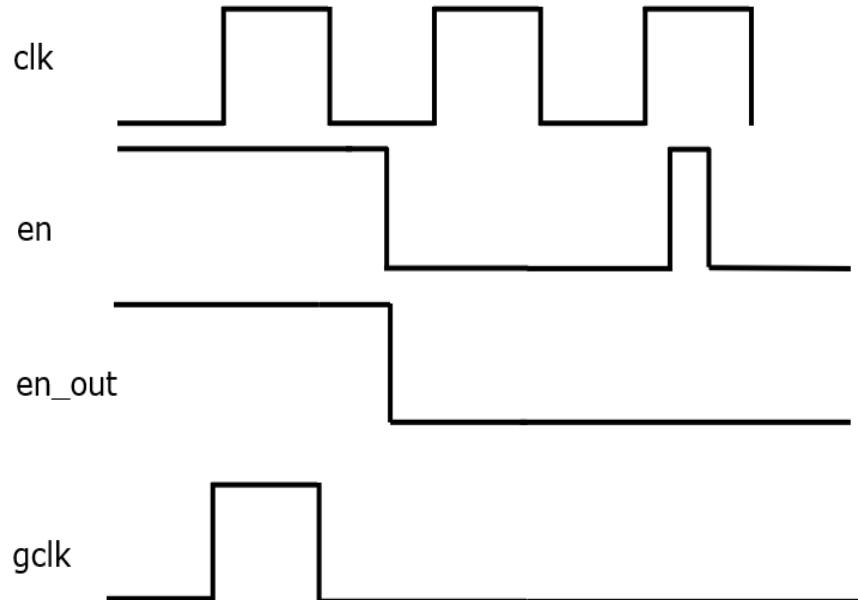
# CLOCK GATING – GLITCH PROBLEM

- What happens if there is a glitch on the enable signal?



# SOLUTION: GLITCH-FREE CLOCK GATE

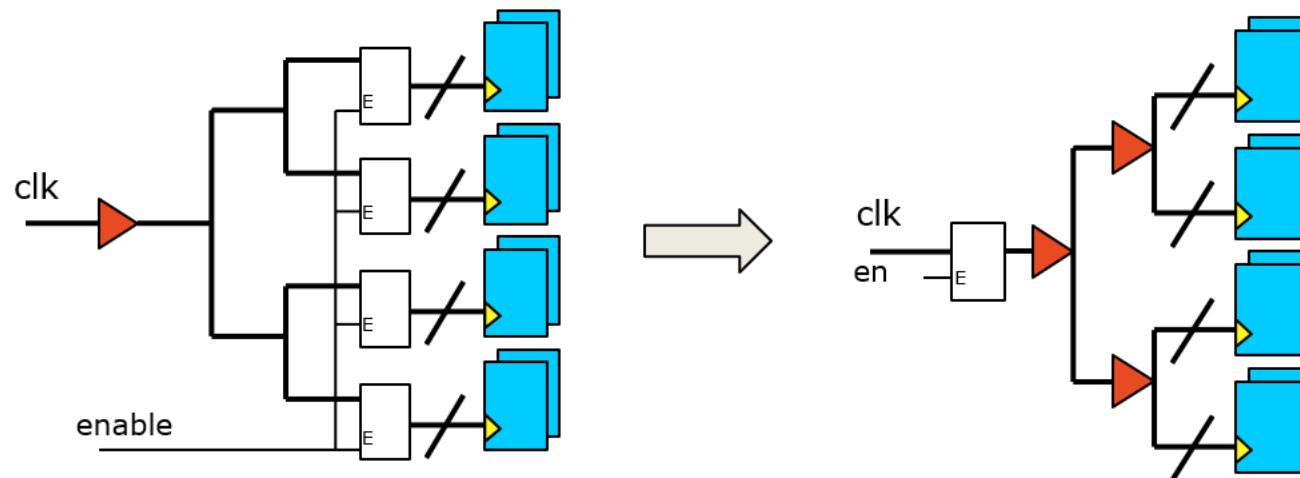
- By latching the enable signal during the positive phase, we can eliminate glitches:



```
//clock gating with glitch prevention latch
always @ (enable or clk)
begin
    if (!clk)
        en_out <= enable;
    end
assign gclk = en_out && clk;
```

# MERGING CLOCK ENABLE GATES

- Clock gates with common enable can be merged
  - Lower clock tree power, fewer gates
  - May impact enable signal timing and skew.

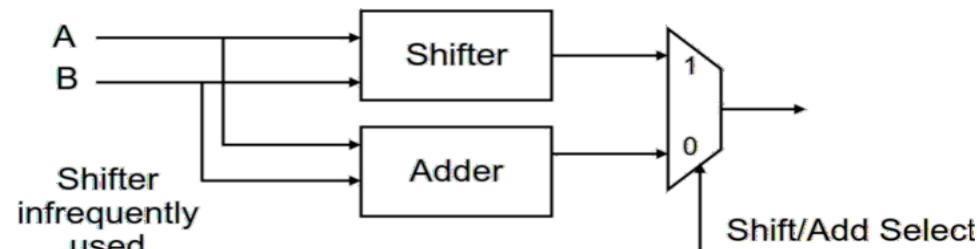


# DATA GATING

- While clock gating is very well understood and automated, a similar situation occurs due to the toggling of data signals that are not used.
- These situations should be recognized and data gated.

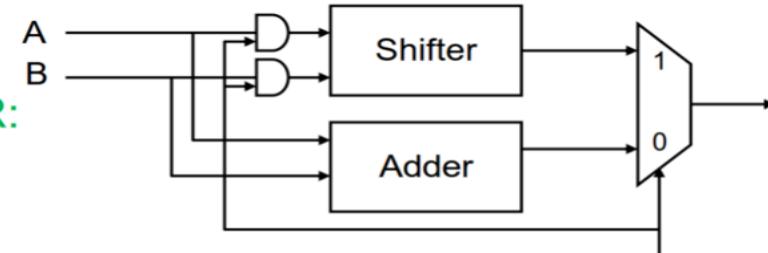
**BAD:**

```
assign add_out = A+B;
assign shift_out = A<<B;
assign out = shift_add ? shift_out : add_out;
```



**BETTER:**

```
assign shift_in_A = A && shift_add;
assign shift_in_B = B && shift_add;
assign shift_out = shift_in_A << shift_in_B;
```



**Thank you!**