

EE 531: ADVANCED VLSI DESIGN

Timing Analysis

Nishith N. Chakraborty

January, 2025


TIMING ANALYSIS

- Static Timing Analysis (STA) – method of computing expected timing of digital circuit without simulation
- Referred to as static because it does not depend on input vectors
- Each cell or module in design is accompanied by timing information (e.g., rise and fall times, delays)


- Timing analysis (usually via STA) must occur between multiple points in the design flow and is used to optimize synthesis, placement, & routing

SYNCHRONOUS DESIGN - REVISIT


- The majority of digital designs are Synchronous and constructed with Sequential Elements.
 - Synchronous design eliminates races.
 - Pipelining increases throughput.
- We will assume that all sequentials are Edge-Triggered, using D-Flip Flops as registers.
- D-Flip Flops have three critical timing parameters:
 - t_{cq} – clock to output: essentially a propagation delay
 - t_{setup} – setup time: the time the data needs to arrive before the clock
 - t_{hold} – hold time: the time the data has to be stable after the clock




TIMING METRICS: RISING EDGE TRIGGERED FLIP-FLOP





TIMING METRICS: RISING EDGE TRIGGERED FLIP-FLOP




TIMING METRICS: RISING EDGE TRIGGERED FLIP-FLOP



TIMING METRICS: RISING EDGE TRIGGERED FLIP-FLOP




TIMING METRICS: RISING EDGE TRIGGERED FLIP-FLOP



TIMING CONSTRAINTS

- There are two main problems that can arise in synchronous logic:
 - Max Delay: The data doesn't have enough time to pass from one register to the next before the next clock edge.
 - Min Delay: The data path is so short that it passes through several registers during the same clock cycle.
- Max delay violations are a result of a slow data path, including the registers' t_{setup} , therefore it is often called the "Setup" path.
- Min delay violations are a result of a short data path, causing the data to change before the t_{hold} has passed, therefore it is often called the "Hold" path.


SETUP TIME CONSTRAINT



SETUP TIME CONSTRAINT



$$T_c < t_{pcq,A} + t_{pd} + t_{\text{SETUPB}}$$

Setup time violation!
Must be avoided



HOLD TIME CONSTRAINT

$$t_{cd} \geq t_{hold} - t_{ccq}$$




TIMING CONSTRAINTS: SUMMARY

- For **Setup** constraints, the data has to propagate fast enough to be captured by **the next clock edge**:
 - This sets our maximum frequency.
 - If we have setup failures, we can always just **slow down the clock**.
- For **Hold** constraints, the data path delay has to be long enough so it isn't accidentally captured by **the same clock edge**:
 - This is independent of clock period.
 - If there is a hold failure, you can throw your chip away!

STATIC TIMING ANALYSIS (STA)

- STA checks the worst case propagation of all possible vectors for min/max delays.
- Advantages:
 - Much faster than timing-driven, gate-level simulation
 - Exhaustive, i.e., **every (constrained) timing path is checked**.
 - Vector generation NOT required




STATIC TIMING ANALYSIS (STA)

- Disadvantages:
 - Proper circuit functionality is NOT checked
 - Must define timing requirements/exceptions
- Limitations:
 - Only useful for synchronous design
 - Cannot analyze combinatorial feedback loops
 - e.g., a flip-flop created out of basic logic gates
 - Cannot analyze asynchronous timing issues
 - Such as clock domain crossing
 - Will not check for glitching effects on asynchronous pins
 - Combinatorial logic driving asynch (set/reset) pins of sequential elements will not be checked for glitching



TIMING PATHS

- A path is a route from a **Startpoint** to an **Endpoint**.
- **Startpoint (SP)**
 - Clock pins of the flip flops
 - Input ports , a.k.a Primary Inputs (PI)
- **Endpoints (EP)**
 - Input pins of the flip flops, except the clock pins
 - Output ports, a.k.a Primary Outputs (PO)
 - Memories / Hard macros
- There can be:
 - Many paths going to any one endpoint
 - Many paths for each start-point and end-point combination



STATIC TIMING ANALYSIS

- Four categories of timing paths
 - Register to Register (**reg2reg**)
 - Register to Output (**reg2out**)
 - Input to Register (**in2reg**)
 - Input to Output (**in2out**)



GOALS OF STATIC TIMING ANALYSIS


- Verify max delay and min delay constraints are met for all paths in a design.
 - Start with a Gate-Level Netlist.
 - Timing Models are provided for every gate in the library.
 - Static Timing Analysis needs to report if any path violates the max/min delay constraints.
- But is this enough?
 - No!
 - We want to know all the paths that violate the timing constraints.
 - In fact, we want to know the timing of all paths reported in order of length.
 - And we want to know where the problems are so we can go about fixing them.
- Let's see the basic idea of how this can be done.

SOME BASIC ASSUMPTIONS


- Our design is synchronous
 - In addition, we will only be showing how to deal with combinational elements and max delay constraints.
- We will assume a pin-to-pin delay model
 - In other words, each gate has a single, constant delay from input to output.
 - In the real world, gate delay is affected by many factors, such as gate type, loading, waveform shape, transition direction, particular pin, and random variation.
 - As we saw earlier, a real design gets all this data from the .lib files.
- We will take a topological approach
 - In other words, we disregard the logical functionality of the gates and therefore, consider all paths, though some of them cannot logically happen.
 - More on this later...

SIMPLE PATH REPRESENTATION


- Let's say we have the circuit:



- And the timing model of our AND gate is:




- Sometimes also described as a combination of input wire delay and gate delay:




SIMPLE PATH REPRESENTATION

- We will build a graph:
 - Vertices: Wires, 1 per gate output and 1 for each SP and EP.
 - Edges: Gates, input pin to output pin, 1 edge per input with a delay for each edge.



- Finally, add Source/Sink Nodes:
 - 0-weight edge to each SP and from each EP.
 - That way all paths start and end at a single node.



NODE ORIENTED TIMING ANALYSIS


- If we would enumerate every path, we would quickly get exponential explosion in the number of paths.
- Instead, we will use node-oriented timing analysis
 - For each node, find the worst delay to the node along any path.
- For this, we need to define two important values:
 - Arrival Time at a node (AT): the longest path from the source to the node.
 - Required Arrival Time at node (RAT): the latest time the signal is allowed to leave the node to make it to the sink in time.
- Slack at node n is defined as: $Slack(n) = RAT(n) - AT(n)$



HOW DO WE COMPUTE ATS AND RATs?

- Recursively!


- The **Arrival Time** at a node is just the maximum of the **ATs** at the **predecessor** nodes **plus** the delay from that node.
- The **Required Arrival Time** to a node is just the minimum of the **RATs** at the **successor** nodes **minus** the delay to that node.




$$AT(n) = \begin{cases} 0 & n = \text{SRC} \\ \max_{p \in \text{pred}(n)} [AT(p) + \Delta(p,n)] & n \neq \text{SRC} \end{cases}$$

$$RAT(n) = \begin{cases} T & n = \text{SNK} \\ \min_{s \in \text{succ}(n)} [RAT(s) - \Delta(n,s)] & n \neq \text{SNK} \end{cases}$$


AT, RAT, AND SLACK: GRAPHICAL OVERVIEW



AT, RAT, AND SLACK: GRAPHICAL OVERVIEW




AT, RAT, AND SLACK: GRAPHICAL OVERVIEW



STA: EXAMPLE


- Just look at this path and try to find the worst path.
 - Does it meet a cycle time of $T=12$?



- Now let's fill in the RAT, AT, and SLACK of each node and:
 - Quickly find out if we meet timing
 - Figure out what the worst path is



STA: EXAMPLE

- We'll start by representing it as a directed acyclic graph (DAG)
- Next, we'll compute ATs from SRC to SNK





STA: EXAMPLE

- And now **RAT** from SNK to SRC




STA: EXAMPLE

- And finally, we can calculate the **slack**.
- And guess what – we found the critical path!



FALSE PATHS


- We saw how to find the RAT, AT and Slack at every node.
 - All of this can be done very efficiently and be adapted for min timing, sequential elements, latch-based timing, etc.
 - Even better, we can quickly report the order of the critical paths.
- However, this was all done topologically (i.e., without looking at logic).
 - Let's see why this is a problem



Design Constraints: Revisited

TIMING CONSTRAINTS

- **Question:** How does the STA tool know what the required clock period is?
- **Answer:**
 - We have to tell it!
 - We have to define constraints for the design.
 - This is usually done using the Synopsys Design Constraints ([SDC](#)) syntax, which is a superset of [TCL](#).
- Three main categories of timing constraints:
 - Clock definitions
 - Modeling the world external to the chip
 - Timing exceptions



TIMING CONSTRAINTS

- EDA tools sometimes use a different data structure called a “collection”
- A collection is similar to a TCL list, but:
 - The value of a collection is not a string, but rather a pointer, and we need to use special functions to access its values.
 - For example, if you were to run `foreach` on a collection, it would just have one element (the pointer to the collection). Instead, use `foreach_in_collection`.
 - I won’t go into the specifics here, but these are some of the collection accessing functions:


`foreach_in_collection`
`index_collection`
`sizeof_collection`
`sort_collection`

`filter_collection`
`add_to_collection`
`compare_collections`

`copy_collection`
`get_object_name`
`remove_from_collection`


DESIGN OBJECTS

- **Design:** A circuit description that performs one or more logical functions (i.e Verilog module).
- **Cell:** An instantiation of a design within another design (i.e Verilog instance).
- **Reference:** The original design that a cell "points to" (i.e Verilog sub-module)
 - Called a module in Stylus Common UI.
- **Port:** The input, output or inout port of a Design.
- **Pin:** The input, output or inout pin of a Cell in the Design.
- **Net:** The wire that connects Ports to Pins and/or Pins to each other.
- **Clock:** Port of a Design or Pin of a Cell explicitly defined as a clock source.




SDC HELPER FUNCTIONS

- Before starting with constraints, let's look at some very useful built in commands:
 - Note that all of these return **collections** and not **TCL lists**!
 - These will only work after design elaboration!
- “get” commands:
 - **[get_ports string]** – returns all ports that match string.
 - **[get_pins string]** – returns all cell/macro pins that match string.
 - **[get_nets string]** – returns all nets that match string.
 - Note that adding the *–hier* option will search hierarchically through the design.



SDC HELPER FUNCTIONS

- “all” commands:
 - [all_inputs] – returns all the primary inputs (ports) of the block.
 - [all_outputs] – returns all the primary outputs (ports) of the block.
 - [all_registers] – returns all the registers in the block.



CLOCK DEFINITIONS

- To start, we must define a clock:
 - Where does the clock come from? (i.e., input port, output of PLL, etc.)
 - What is the clock period? (=operating frequency)
 - What is the duty-cycle of the clock?

```
create_clock -period 20 -name my_clock [get_ports clk]
```

- Can there be more than one clock in a design?
 - Yes, but be careful about clock domain crossings!
 - If a clock is produced by a clock divider, define a “[generated clock](#)”:

```
create_generated_clock -name gen_clock \
                      -source [get_ports clk] -divide_by 2 [get_pins FF1/Q]
```

CLOCK DEFINITIONS

- But during synthesis, we assume the clock is ideal, so:

```
set_ideal_network [get_ports clk]
```

- However, for realistic timing, it should have some transition:

```
set_clock_transition 0.2 [get_clocks my_clock]
```

- And we may want to add some jitter, so:

```
set_clock_uncertainty 0.2 [get_clocks my_clock]
```


- Finally, after building a clock tree, we do not want the clock to be ideal anymore, so:

```
set_propagated_clock [get_clocks my_clock]
```

I/O CONSTRAINTS

- Now that the clock is defined, reg2reg paths are sufficiently constrained.
- However, what about in2reg, reg2out, and in2out paths?
 - First, what clock toggles an I/O port?
 - And what about the time needed outside the chip?
- Define I/O constraints:
- Input and output delays model the length of the path outside the chip:

```
set_input_delay 0.8 -clock clk \
    [remove_from_collection [all_inputs] [get_ports clk]]
set_output_delay 2.5 -clock clk [all_outputs]
```



I/O CONSTRAINTS

- An alternative approach is to define max delays to/from I/Os:

```
set_max_delay 5 \
    -from [remove_from_collection [all_inputs] [get_ports clk]]
set_max_delay 5 -to [all_outputs]
```

- Additionally, we must model the transitions on the inputs:


```
set_driving_cell -cell [get_lib_cells MYLIB/INV4] -pin Z \
    [remove_from_collection [all_inputs] [get_ports clk]]
```

- And capacitance of the outputs:

```
set_load $CIN_OF_INV [all_outputs]
```


I/O CONSTRAINTS

- Graphically, we can summarize the I/O constraints, as follows:



TIMING EXCEPTIONS

- There are several cases when we need to define exceptions that should be treated differently by STA.
- For example, looking into the topology of the network we saw earlier:




- In this case, we would define a **false path**:

```
set_false_path -through [get_pins mux1/I0] -through [get_pins mux2/I0]  
set_false_path -through [get_pins mux1/I1] -through [get_pins mux2/I1]
```

TIMING EXCEPTIONS

- Another common case of a false path is a clock domain crossing through a synchronizer:




```
set_false_path -from F1/CP -to F2/D
```

- Alternatively, this can be defined with:

```
set_clock_groups -logically_exclusive \
    -group [get_clocks C1] -group [get_clocks C2]
```

TIMING EXCEPTIONS

- If an equal-phase (divided) slow clock is sending data to a faster clock, a multi-cycle path may be appropriate:



```
set_multicycle_path -setup -from F1/CP -to F2/D      2  
set_multicycle_path -hold      -from F1/CP -to F2/D    1
```

TIMING EXCEPTIONS

- A common case for designs is that some value should be assumed constant
 - For example, setting a register for a certain operating mode.
- In such cases, many timing paths are false
 - For example, if the constant sets a multiplexer selector.
 - Or a '0' is driven to one of the inputs of an AND gate.
- To propagate these constants through the design and disable irrelevant timing arcs, a [set_case_analysis](#) constraint is used:

```
set_case_analysis 0 [get_ports TEST_MODE]
```

DESIGN RULE VIOLATIONS (DRV)

- You can set specific design rules that should be met, for example:
 - Maximum transition through a net.

```
set_max_transition $MAX_TRAN_IN_NS
```

- Maximum Capacitive load of a net.

```
set_max_capacitance $MAX_CAP_IN_PF
```

- Maximum fanout of a gate.

```
set_max_fanout $MAX_FANOUT
```

Thank you!